



Real-Time AI & Agentic Systems with Redpanda

Harnessing Agentic AI & Redpanda to Automate, Optimize, and Scale AI-Driven Workflows

ABOUT ME



DR. HITESH KUMAR SHARMA

*Technical Instructor &
Consultant*

- Industry/Academic Experience: **16 Years (6 years Exp Kafka, Confluent, Redpanda, Python Development & Testing, 6 years in DevOps, K8S & Docker, 4 Years in RPA)**
- Official Instructor of Confluent and Apache Kafka, Redpanda (**50+ Trainings delivered**)
- Delivered **120+ Trainings (LinkedIn, Snapchat, Walmart, HCL, Accenture etc.)**
- Delivered training to **30+ Companies**
- Delivered training in **12+ Countries physically**
- Core Technical Domains: **Java, Test Automation, Confluent, UiPath RPA, DevOps, Cloud Computing, Data Analytics**
- Academic Qualifications: **Ph.D. (CSE), M.Tech (CSE)**
- Certifications:
 - **Confluent Developer, Admin Certified**
 - **UiPath RPA Certified Associate**
 - **Docker Certified Associate**
 - **CKA Certified**
 - **Neo4J Certified Associate**
 - **Maven Certified Professional**
- **5 Books Published**
- **45 Patents Published**
- **02 Copyright Published**



Redpanda Platform & Operations (DevOps/Platform Focus)

Redpanda Fundamentals



- What is Redpanda?
- Broker architecture and IDs
- Kafka basics: topics, partitions, producers/consumers, offsets
- Redpanda vs Apache Kafka performance insights
- Raft consensus algorithm
- Single binary design, thread-per-core model
- Performance metrics on hardware

What is Redpanda?

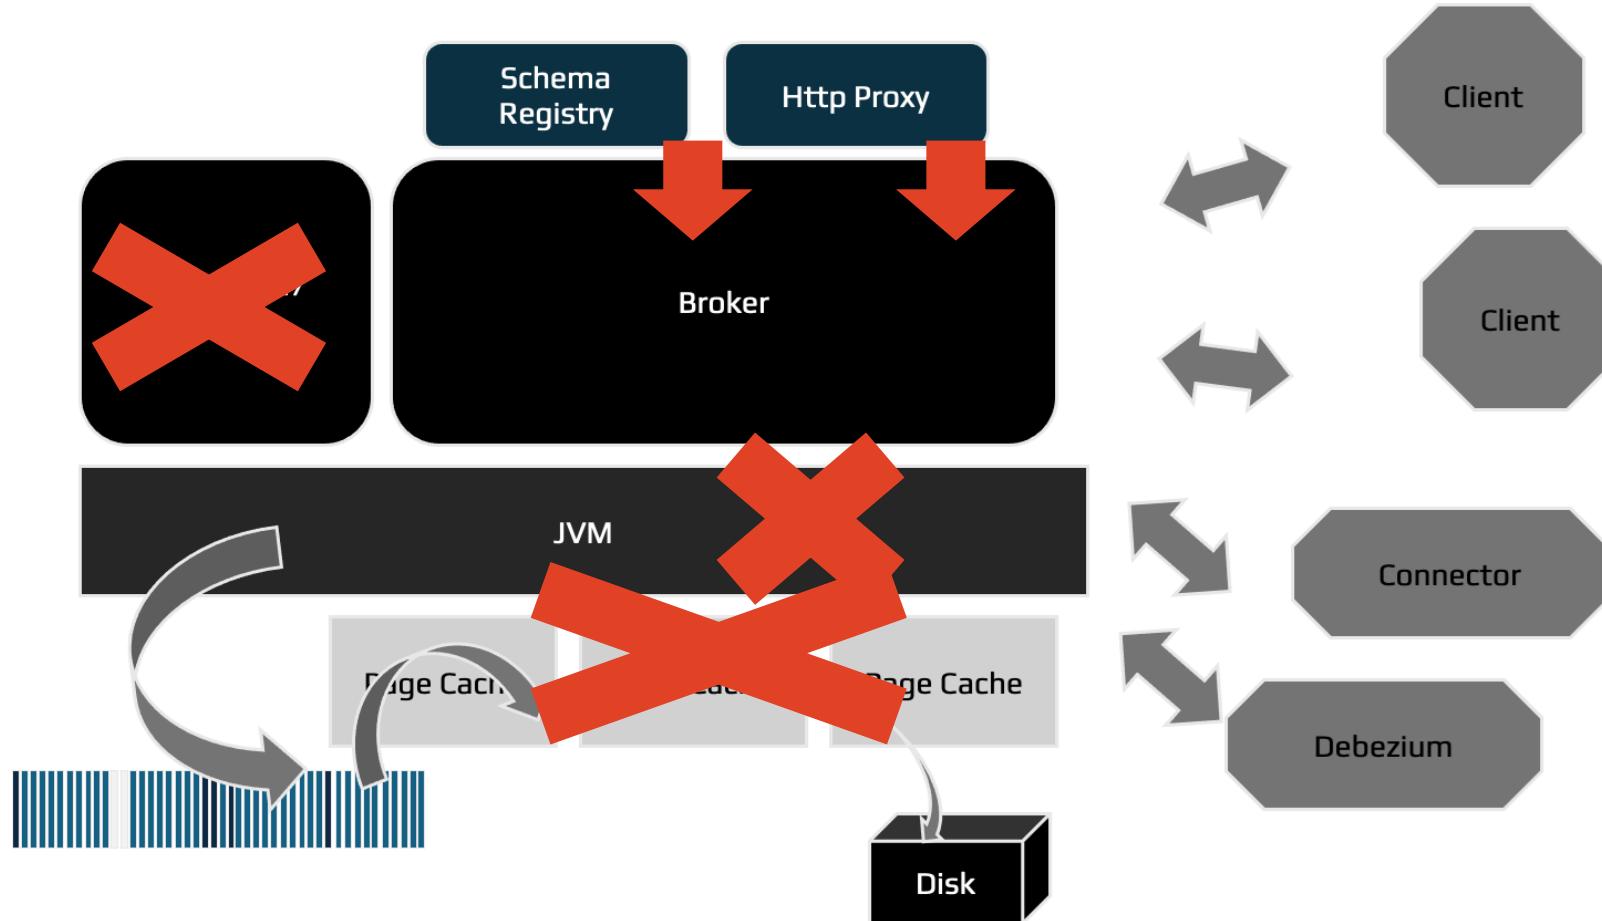
Redpanda is a high-performance streaming data platform, fully compatible with Kafka APIs but simpler to run, faster, and does not require ZooKeeper.

- Low latency and high throughput.
- Kafka API compatible (no code changes needed).
- Simplified architecture (no ZooKeeper).
- Runs anywhere - on-premises, cloud, or containers.



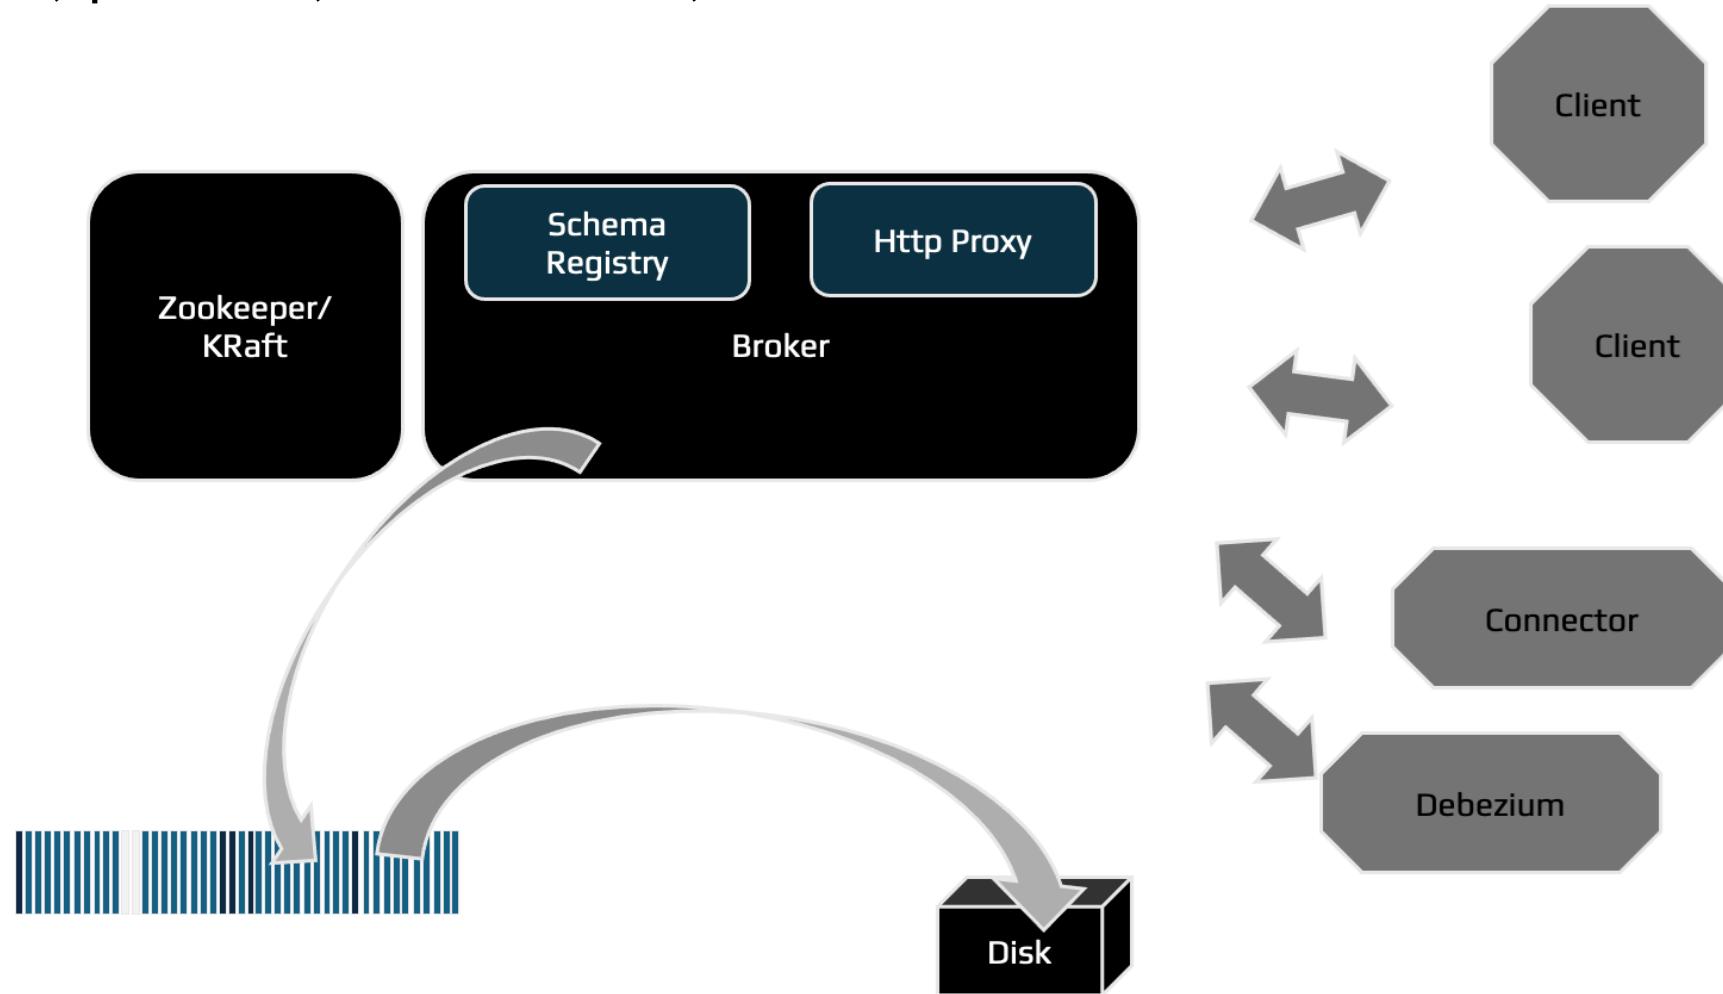
Redpanda Broker architecture

Simple, powerful, cost-efficient, reliable



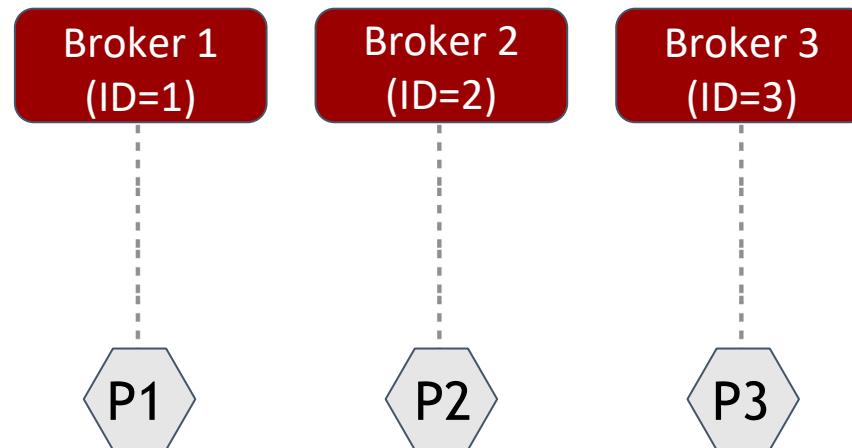
Redpanda Broker architecture

Simple, powerful, cost-efficient, reliable



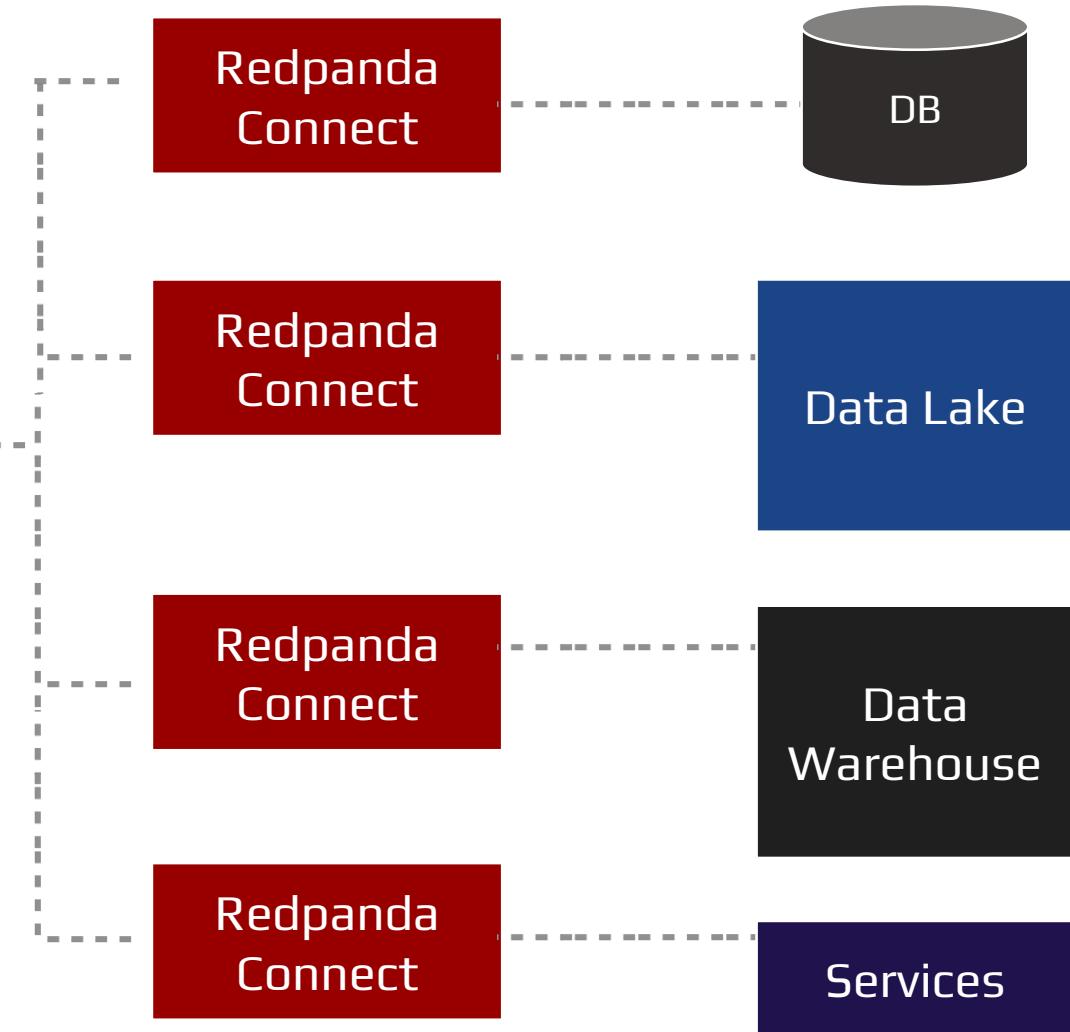
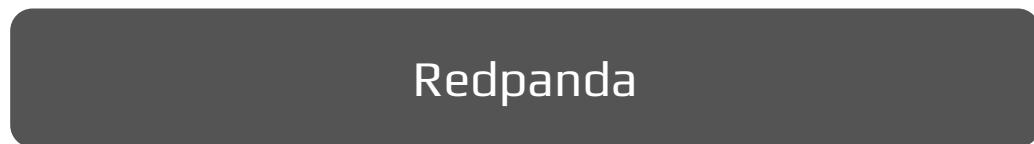
Redpanda Broker IDs

- A **Broker ID** is a **unique identifier** assigned to each broker (server/node) in the Redpanda cluster.
- It helps the cluster **distinguish one broker from another**.
- Even if a broker restarts, the same ID ensures the cluster recognizes it as the same node.

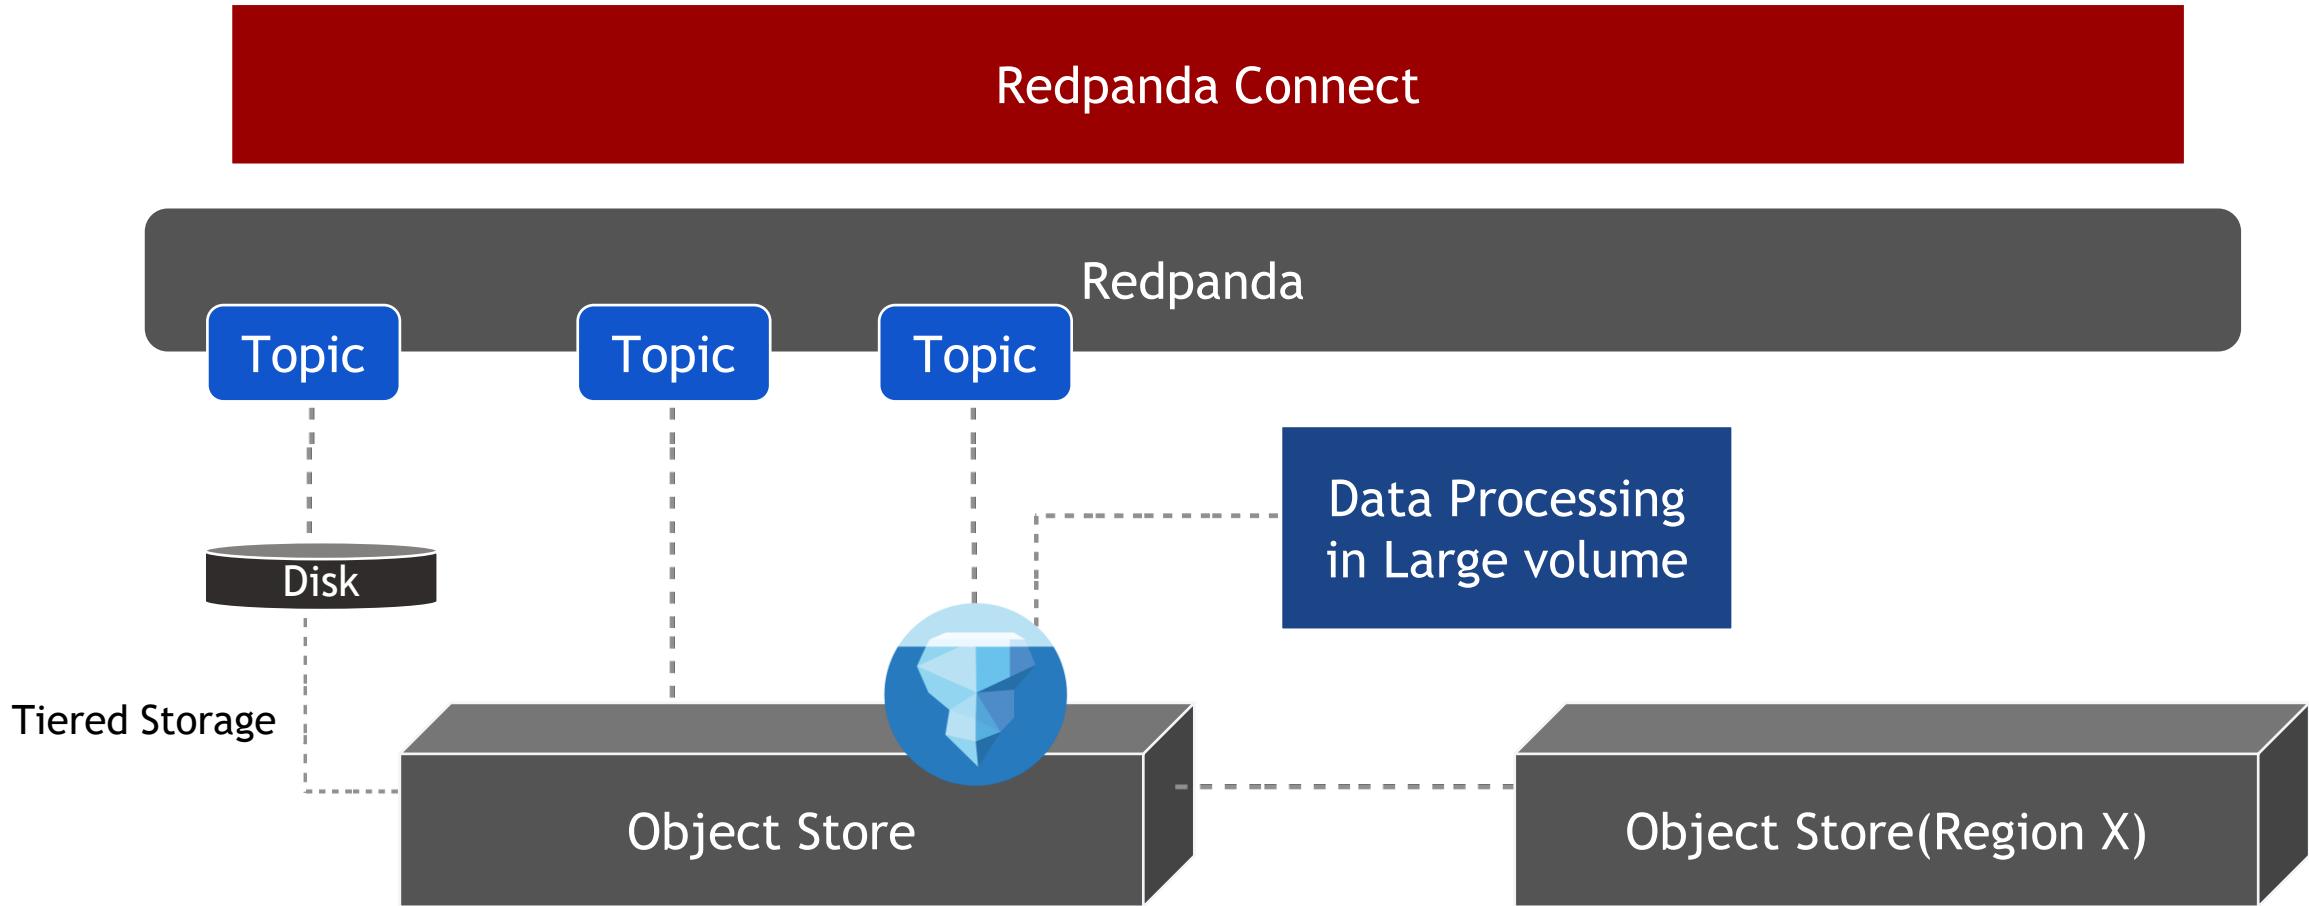


Redpanda Connect

Redpanda Connect is a framework for moving data in and out of Redpanda without writing custom code.



Redpanda One



Kafka basics



Topics

Logical channels where messages are published.

Partitions

Topics are split into partitions for scalability and parallelism.

Producers

Applications that publish (write) data to topics.

Consumers

Applications that subscribe (read) data from topics.

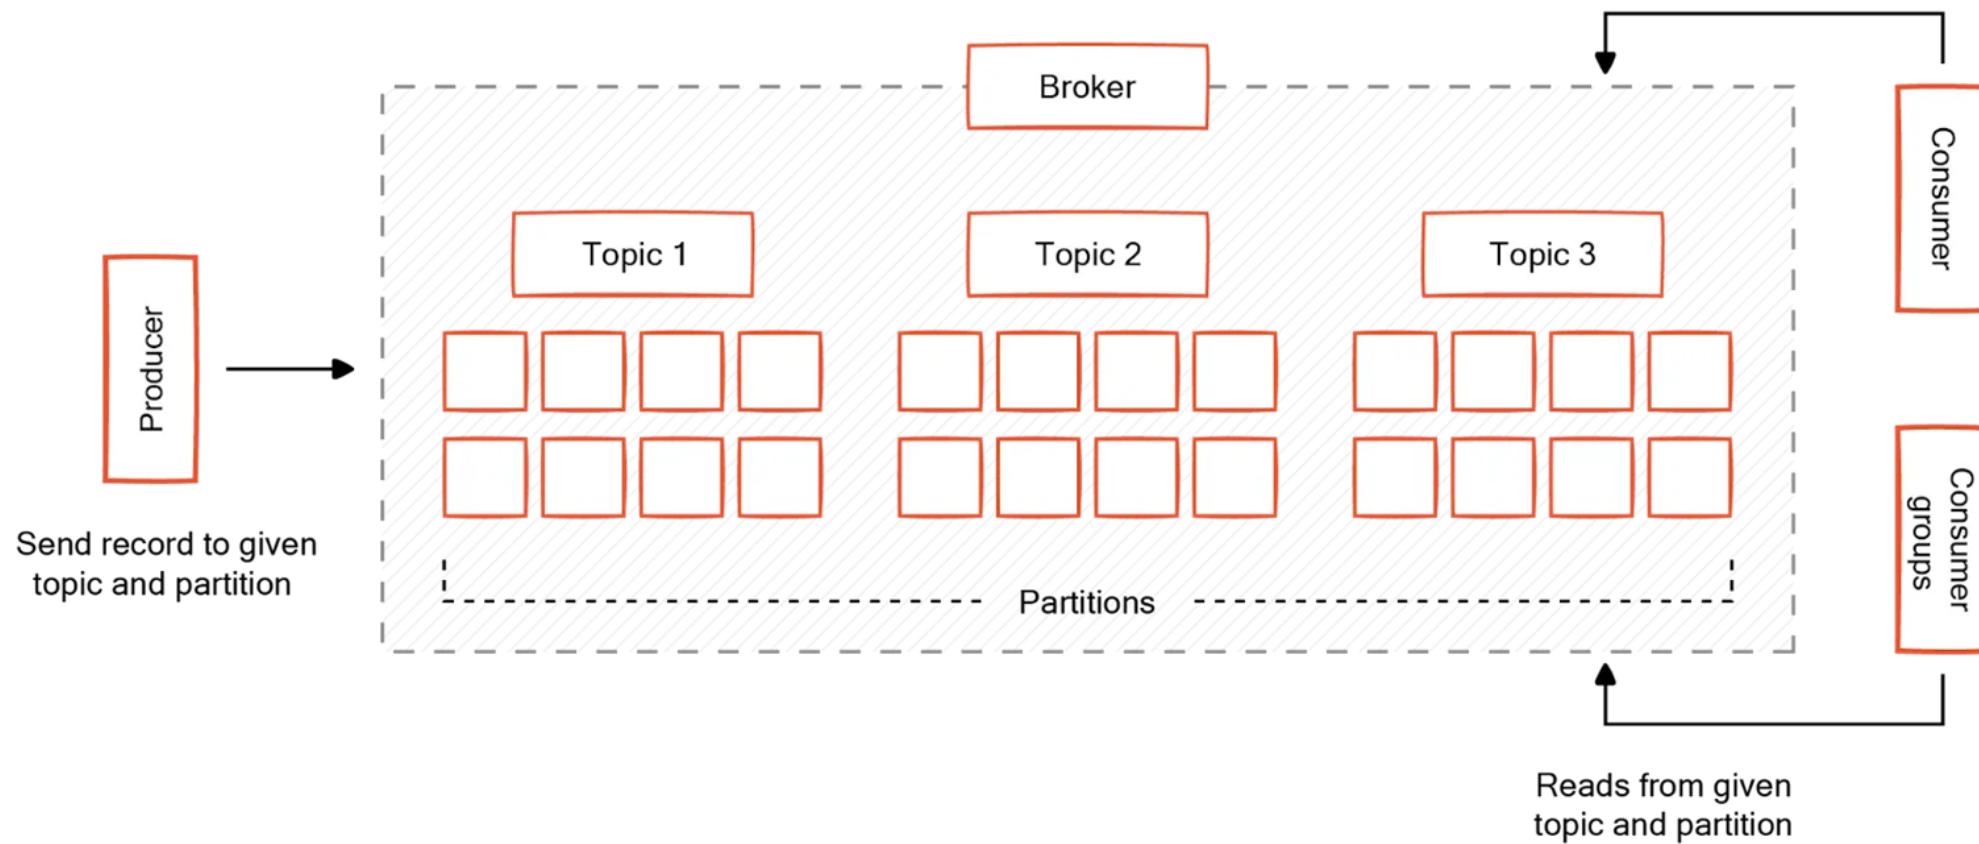
Offsets

Unique sequence numbers within partitions that track message position.

Kafka Topics



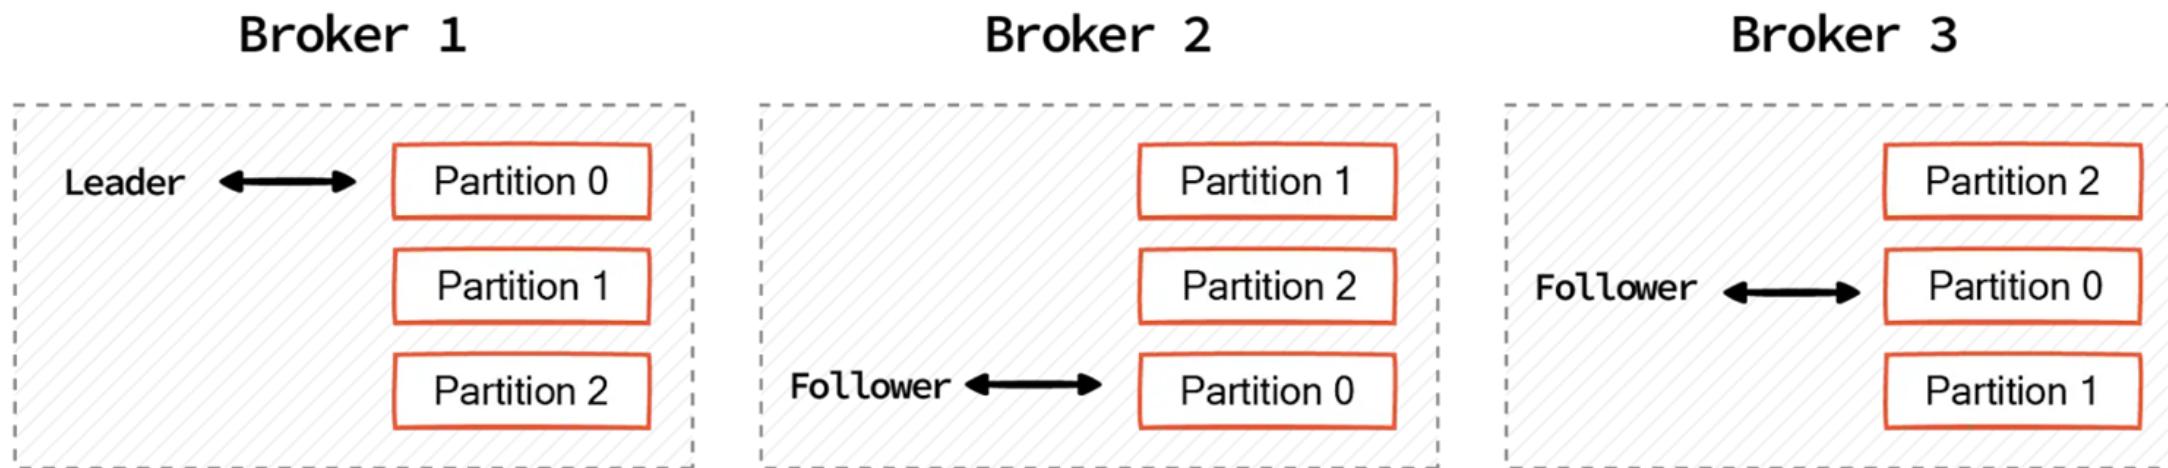
Partitions are ordered, append-only logs that store subsets of a topic's data. Multiple partitions allow Kafka to process messages in parallel.



Kafka Partitions



Partitions are ordered, append-only logs that store subsets of a topic's data. Multiple partitions allow Kafka to process messages in parallel.

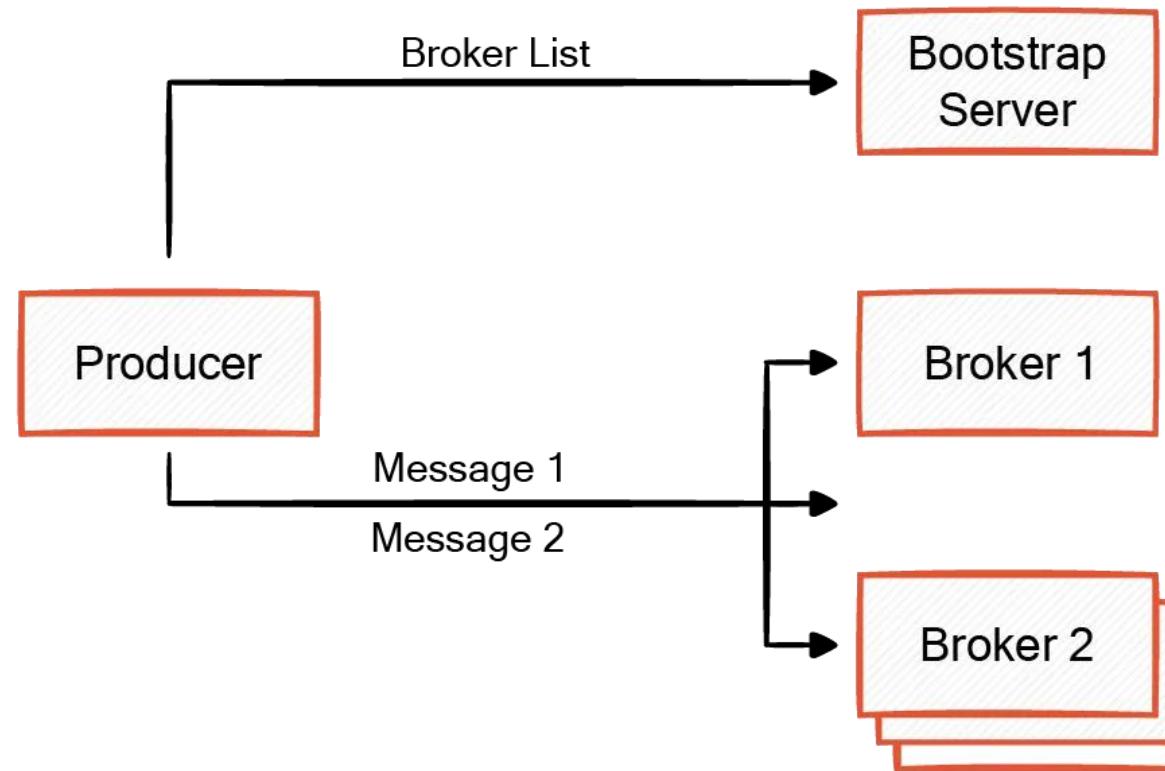


Kafka producers/consumers



When processing data using Apache Kafka's architecture, components connected to Kafka are producers or consumers.

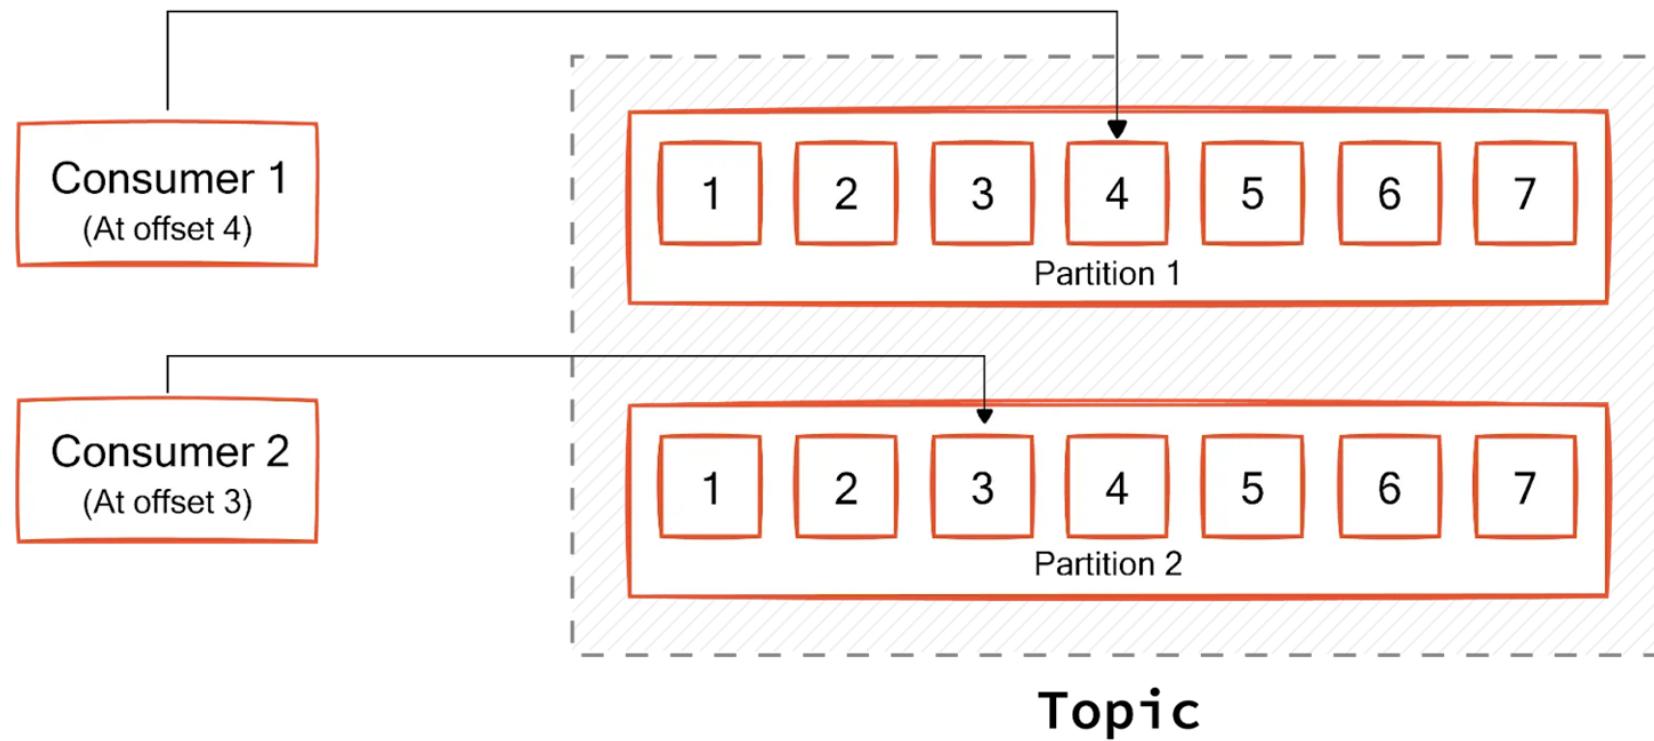
Producers store messages (or events) in Kafka, and consumers read these messages.



Kafka Offset



Offset is a continuously increasing identifier that represents the order of a message from the beginning of the partition.



Redpanda vs Apache Kafka performance insights

Redpanda

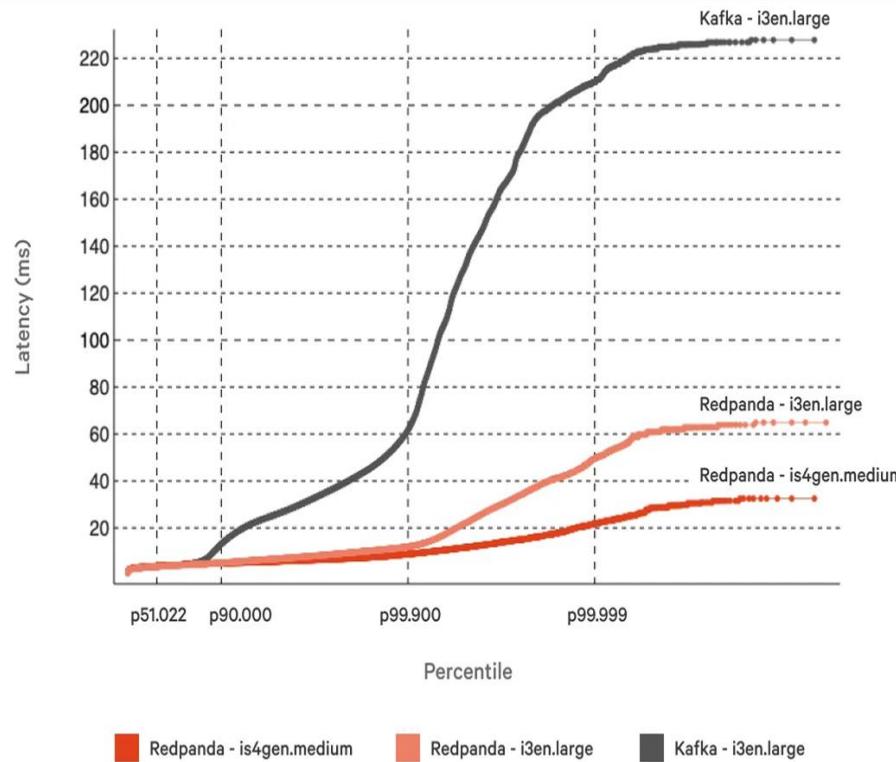
Ultra-low, stable latency (up to 70× faster), breaks 1 GB/s on just 3 nodes, efficient even on small ARM instances.

Apache Kafka

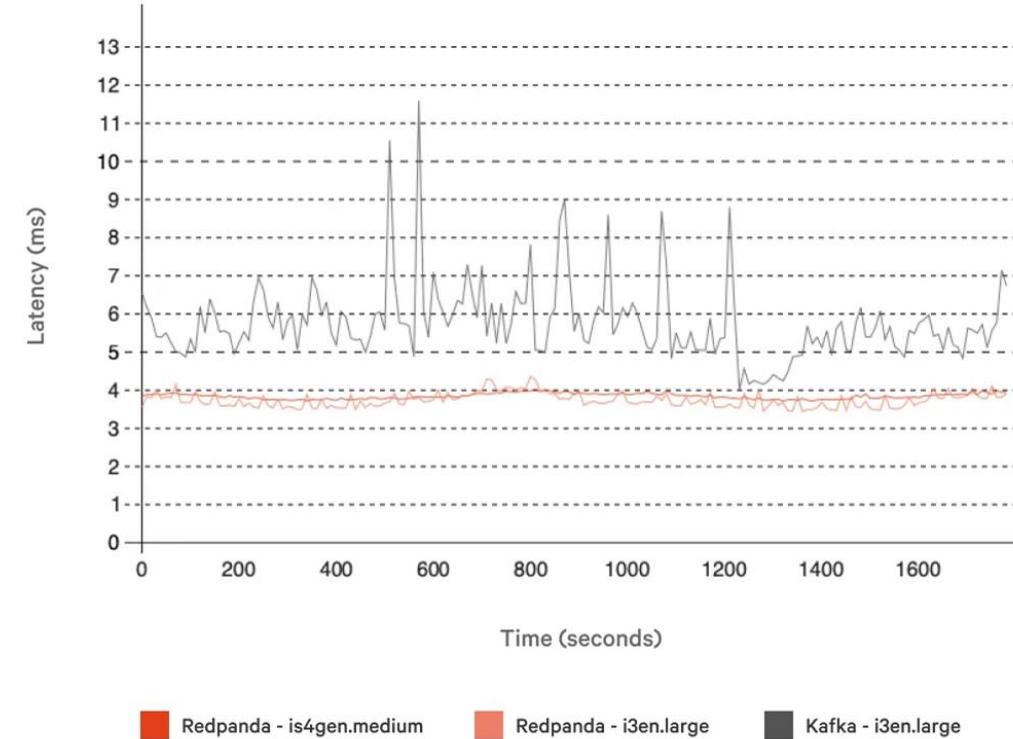
Higher latency, needs 2–3× more nodes and hardware to reach similar throughput, struggles on ARM.

Workload size/Throughput	Kafka P99.99	Redpanda P99.99	Redpanda faster by
Small - 50MB/s	164.57ms (3 nodes)	13.91ms (3 nodes)	12x
Medium - 500MB/s	388.66ms (4 nodes)	16.658ms (3 nodes)	23x
Large - 1GB/s	5509.73ms (6 nodes)	79.574ms (3 nodes)	70x

Comparing performance at 50MB/sec

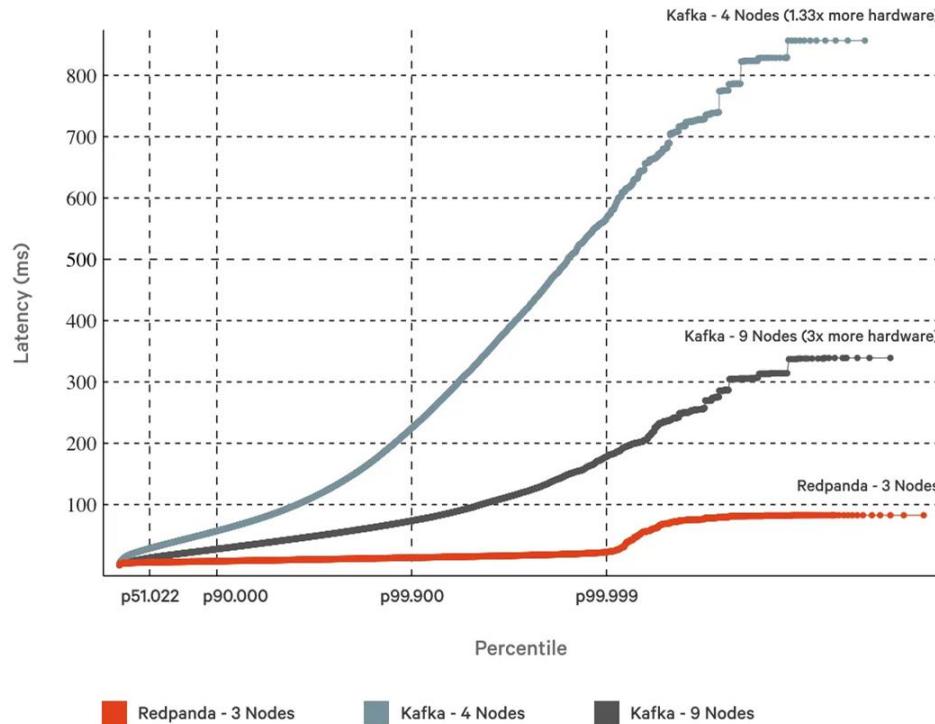


50MB/sec: End-to-End Latency
Percentiles: lower is better

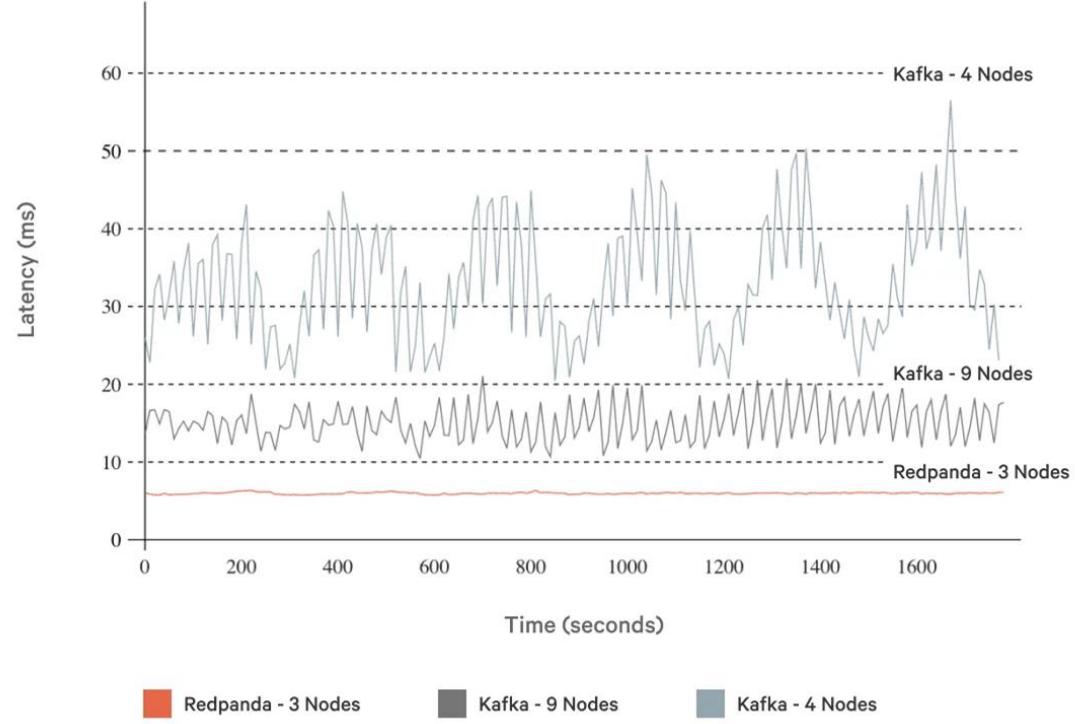


50MB/sec: End-to-End Latency
Average: lower is better

Comparing performance at 500MB/sec

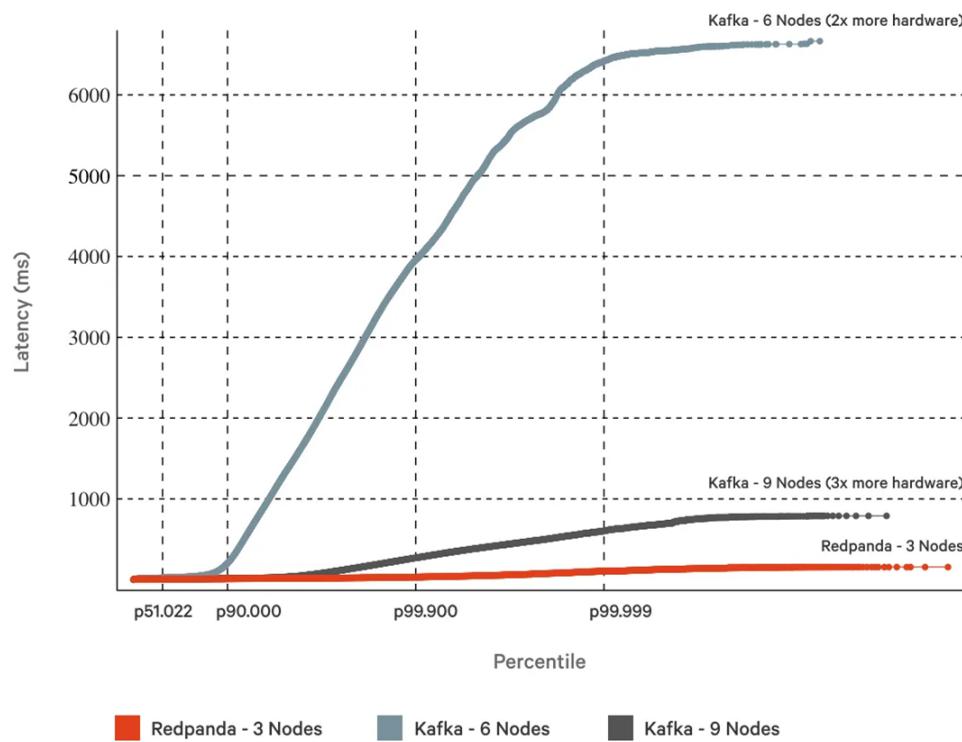


500MB/sec: End-to-End Latency
Percentiles: lower is better

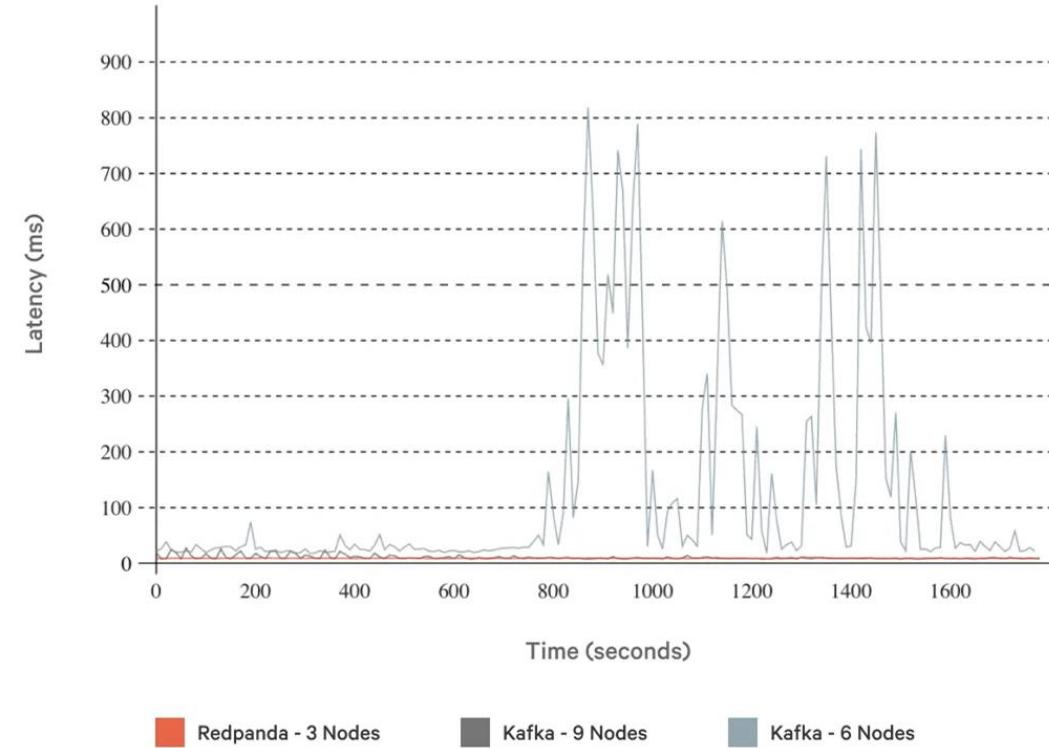


500MB/sec: End-to-End Latency
Average: lower is better

Comparing performance at 1GB/sec



1GB/sec: End-to-End Latency
Percentiles: lower is better



1GB/sec: End-to-End Latency
Average: lower is better

Raft consensus algorithm

Redpanda uses the well-known Raft distributed consensus algorithm as a foundation for the distributed log.

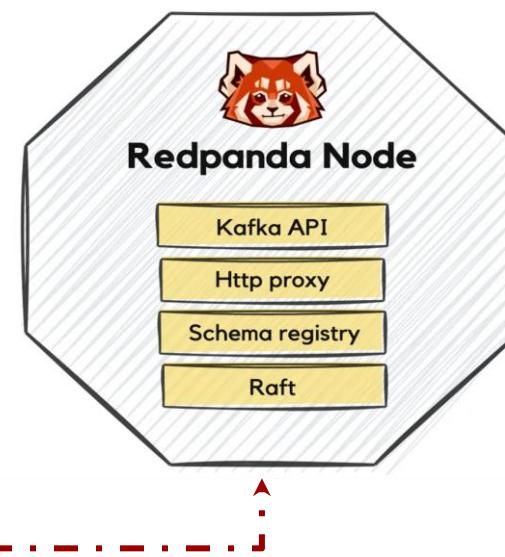
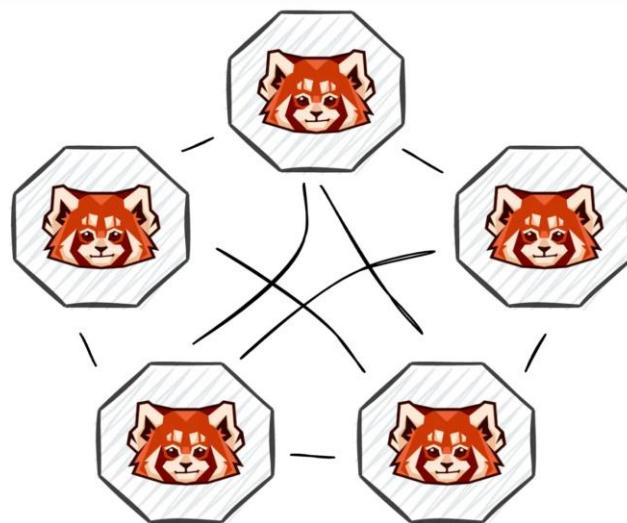
Some of the custom extensions that exist in Redpanda Raft are:

- priority based voting
- prevotes
- reconfiguration
- learners support
- out of band heartbeat generation
- flush debouncing
- support for Apache Kafka®'s ACKS



Single binary design

- Redpanda is distributed as a single executable binary.
- Unlike Kafka, which needs JVM, ZooKeeper, and multiple configs, Redpanda packages everything in one binary.



Thread-per-core model

Redpanda uses a thread-per-core architecture

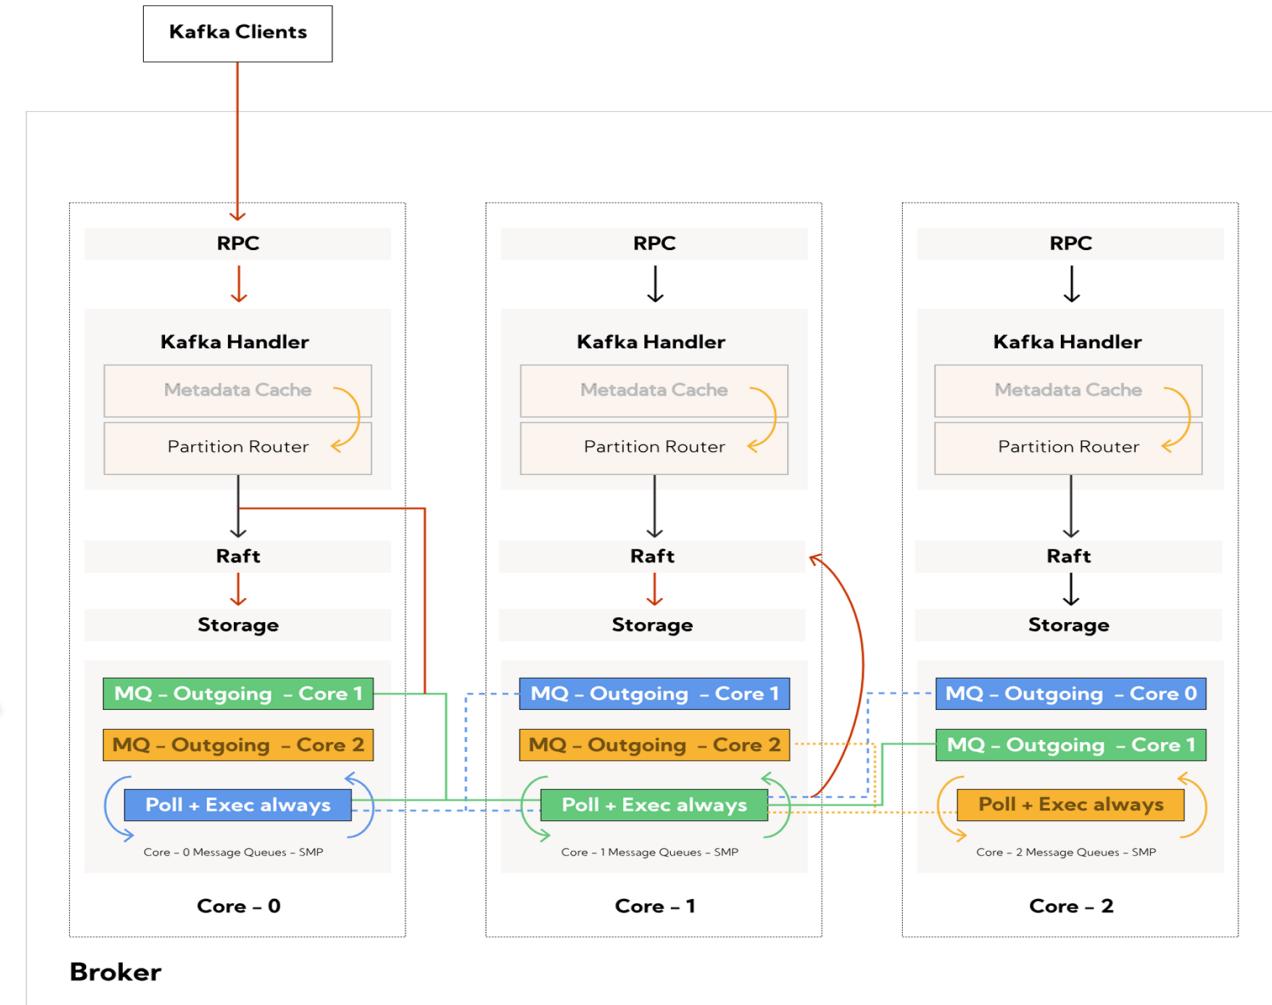
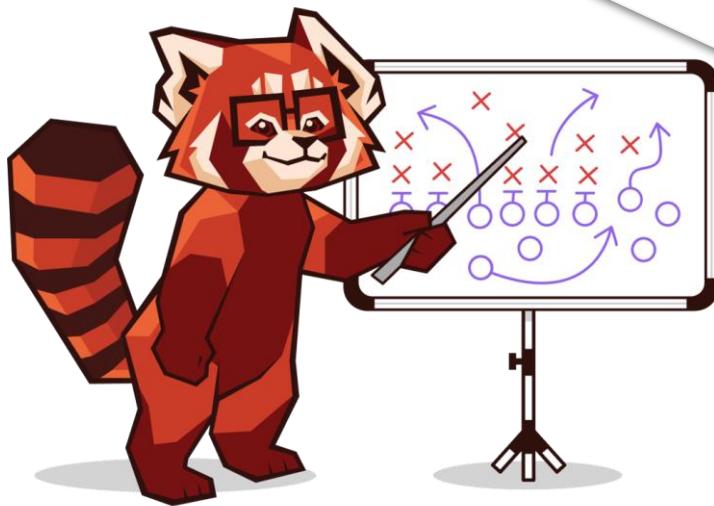
Instead of many threads competing across CPUs , each CPU core runs its dedicated thread, These threads communicate via message passing instead of locks.

Benefits:

- Ultra-low latency** → avoids lock contention and CPU bottlenecks.
- Predictable performance** → each core works independently.
- High throughput** → can process millions of messages/sec efficiently.

Request flow architecture

Redpanda uses a single pinned thread per core to handle all tasks—network polling, async I/O, events, timers, and compute—ensuring no task blocks longer than 500 microseconds to avoid added latency.



Performance metrics on hardware

In Redpanda, the C++ engine and thread-per-core model maximize hardware utilization for higher throughput, lower latency, and better efficiency.

Why Hardware Efficiency Matters:

- Redpanda can deliver 4-10x better performance than Kafka on the same hardware.
- Reduces operational cost by doing more with fewer brokers.
- Well-suited for latency-sensitive applications (trading systems, IoT, real-time analytics).



Redpanda Platform & Operations (DevOps/Platform Focus)

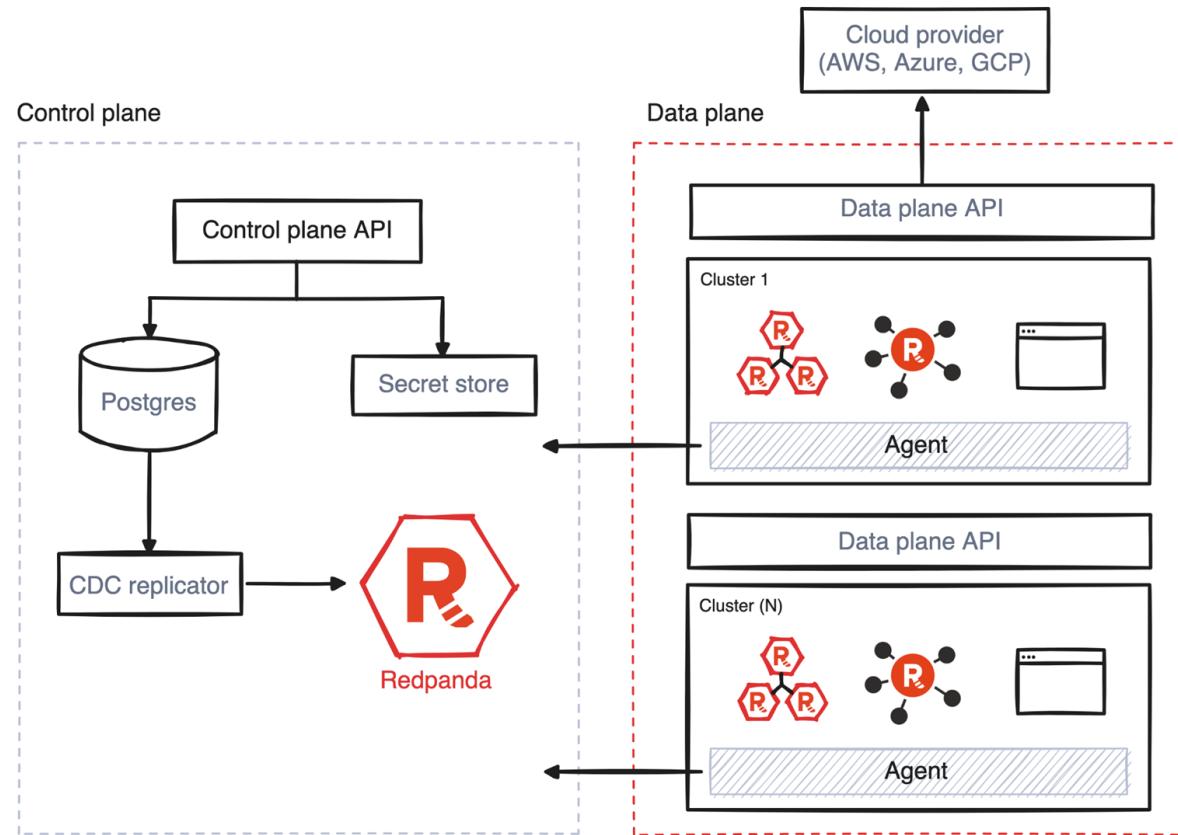
Resilience Concepts in a BYOC Deployment



- Replication, partitions, compaction overview
- Designing for durability and high availability in managed context
- Monitoring, SLAs, Dead-Letter Queues, retry handling
- Observability patterns post-MSK lessons

Redpanda BYOC

In BYOC, the control plane manages clusters, policies, and RBAC, while the data plane runs in your VPC to handle topics, consumer groups, connectors, and schemas.



Why Resilience Matters

Real-world outages: finance transactions dropped, compliance gaps.



Even if outages happen in the real world, using Redpanda in a BYOC setup ensures you still get enterprise-grade resilience while keeping control over your cloud infrastructure.

Core Building Blocks

Replication:

Data copied across brokers for durability

Partitions:

Enable parallelism + fault isolation

Log Compaction:

Keeps latest state while saving storage

High Availability:

Cluster-level fault tolerance

Mini quiz:

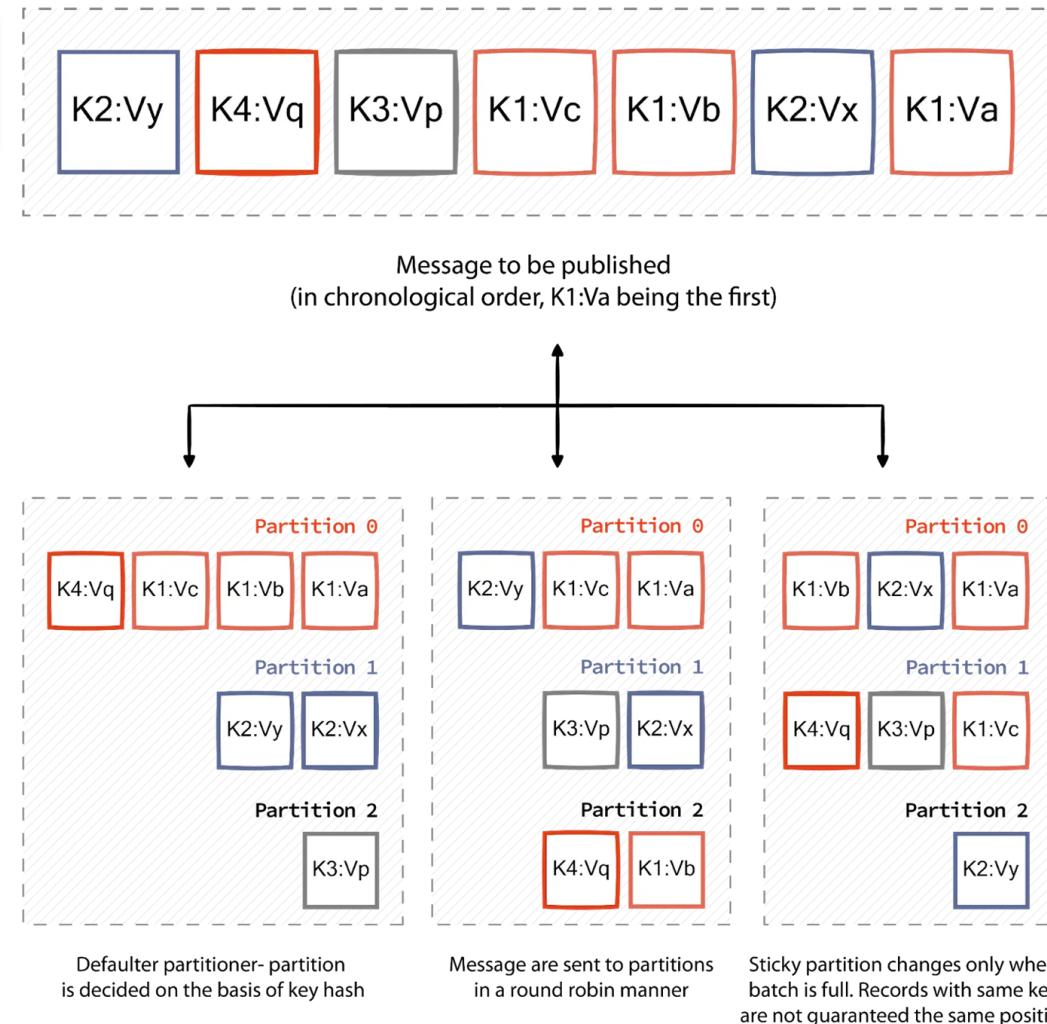
“If RF=2 and 1 broker fails, do we lose data?”

Consumers and Kafka partitions

Kafka guarantees in-order delivery within a partition, and producer partitioning strategies decide which partition a message is written to.

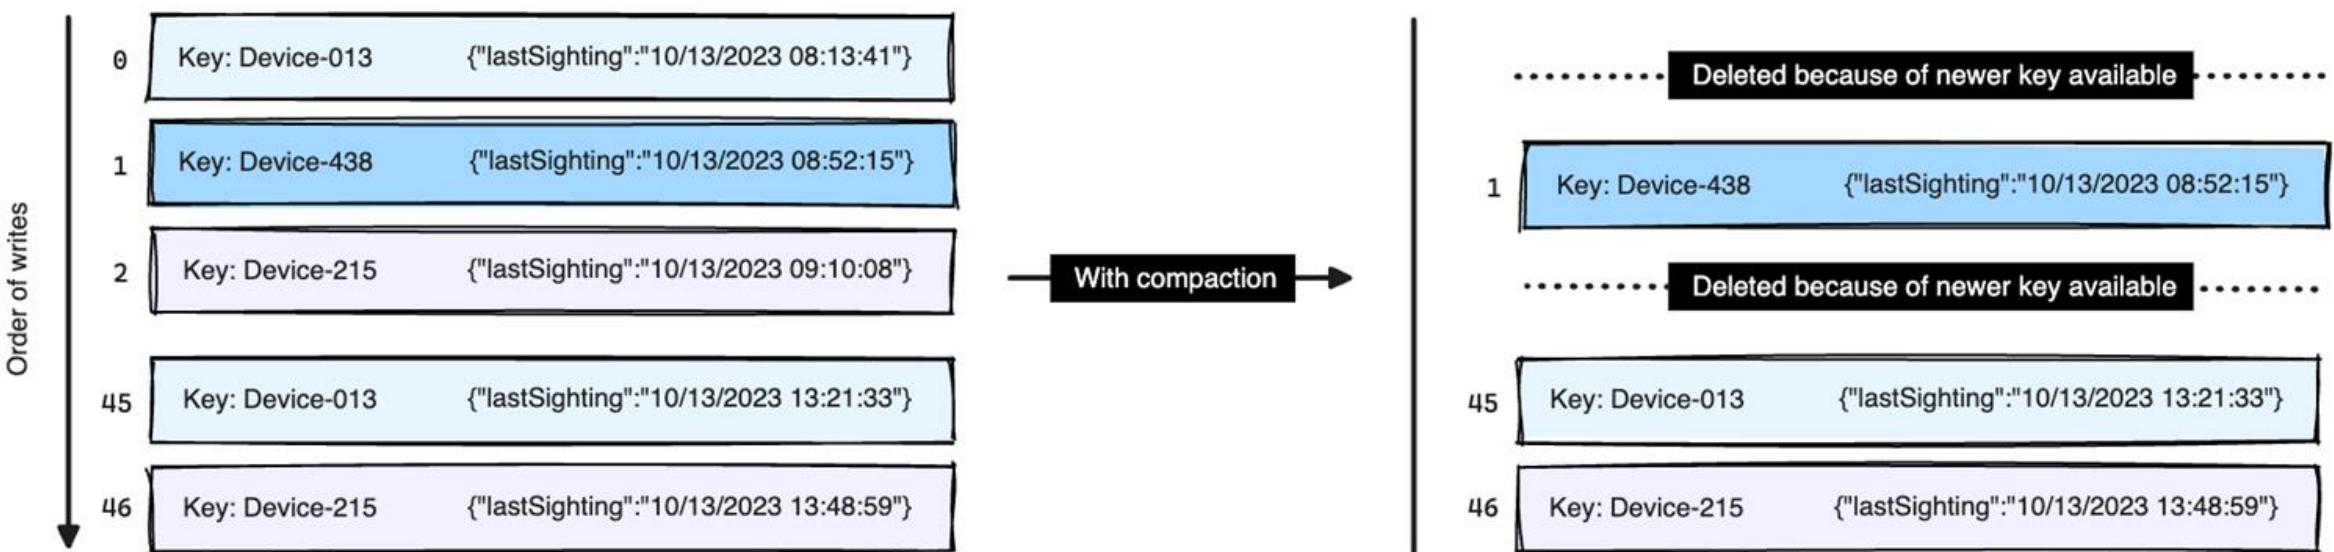
Kafka example message:

- **Key:** "Temperature sensor"
- **Value:** "Recorded a temperature of 25°C"
- **Timestamp:** "April 26, 2024 at 1:00 p.m."



Redpanda compaction overview

Compaction in Redpanda saves storage by keeping only the latest record per key and deleting older duplicates, ensuring topics reflect the most recent state.



Durability and high availability in managed context

Redpanda is designed to ensure data integrity and high availability (HA), even at high-throughput levels.

Deployment strategies:

Failure	Impact	Mitigation strategy
Broker failure	Loss of function for an individual <u>broker</u> or for any virtual machine (VM) that hosts the broker	Multi-broker deployment
Rack or switch failure	Loss of brokers/VMs hosted within that <u>rack</u> , or loss of connectivity to them	Multi-broker deployment spread across multiple racks or network failure domains
Data center failure	Loss of brokers/VMs hosted within that data center, or loss of connectivity to them	Multi-AZ or replicated deployment
Region failure	Loss of brokers/VMs hosted within that region, or loss of connectivity to them	Geo-stretch (latency dependent) or replicated deployment

HA features in Redpanda

Replica Synchronization

Ensures data consistency using Raft quorum-based replication.

Rack Awareness

Distributes replicas across racks or zones to avoid single points of failure.

Partition Leadership

Raft heartbeat and leader elections maintain availability during failures.

Producer Acknowledgment

Acks settings define durability vs latency trade-offs in message delivery.

Partition Rebalancing

Automatically redistributes partitions when brokers are added, removed, or fail.

Tiered Storage & Disaster Recovery

Provides long-term retention and fast recovery of data after failures.

Monitor Redpanda

Redpanda exports metrics through two endpoints on the Admin API port (default: 9644) for you to monitor system health and optimize system performance.

Metrics Endpoints in Redpanda:

Metrics (Internal Metrics)

- Legacy endpoint with many low-level internal metrics
- Best for development, testing, and deep analysis

Public_metrics (Recommended)

- Smaller, curated set of key metrics
- Faster to query and cheaper to ingest
- Suitable for production monitoring

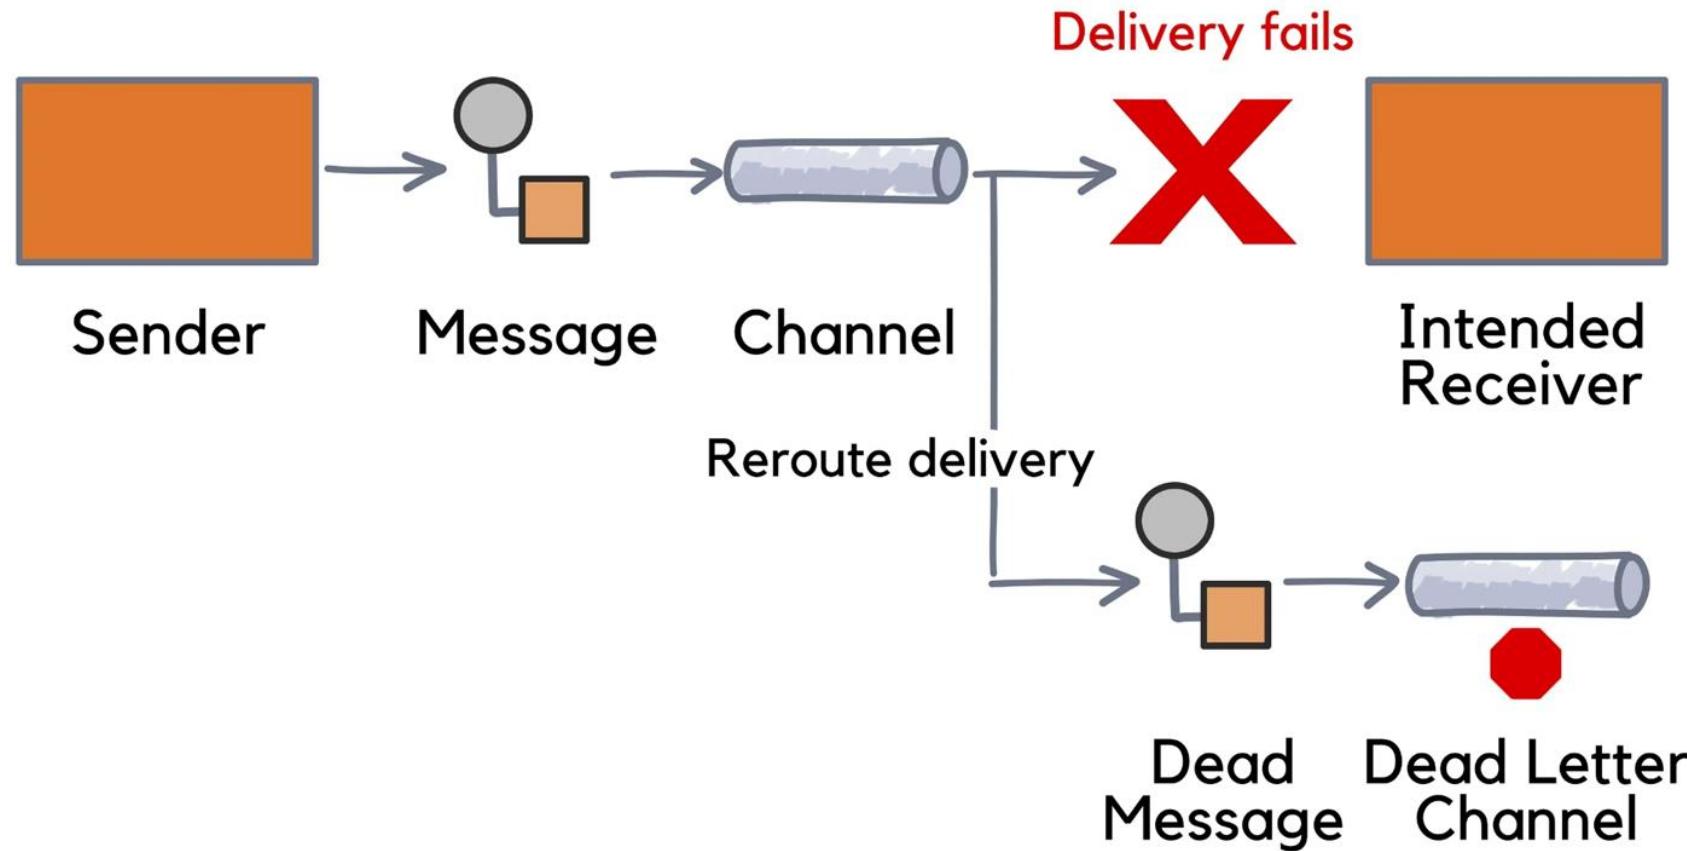
Use Redpanda monitoring

Redpanda provides a repository with examples of monitoring Redpanda with Prometheus and Grafana: [redpanda-data/observability](#)



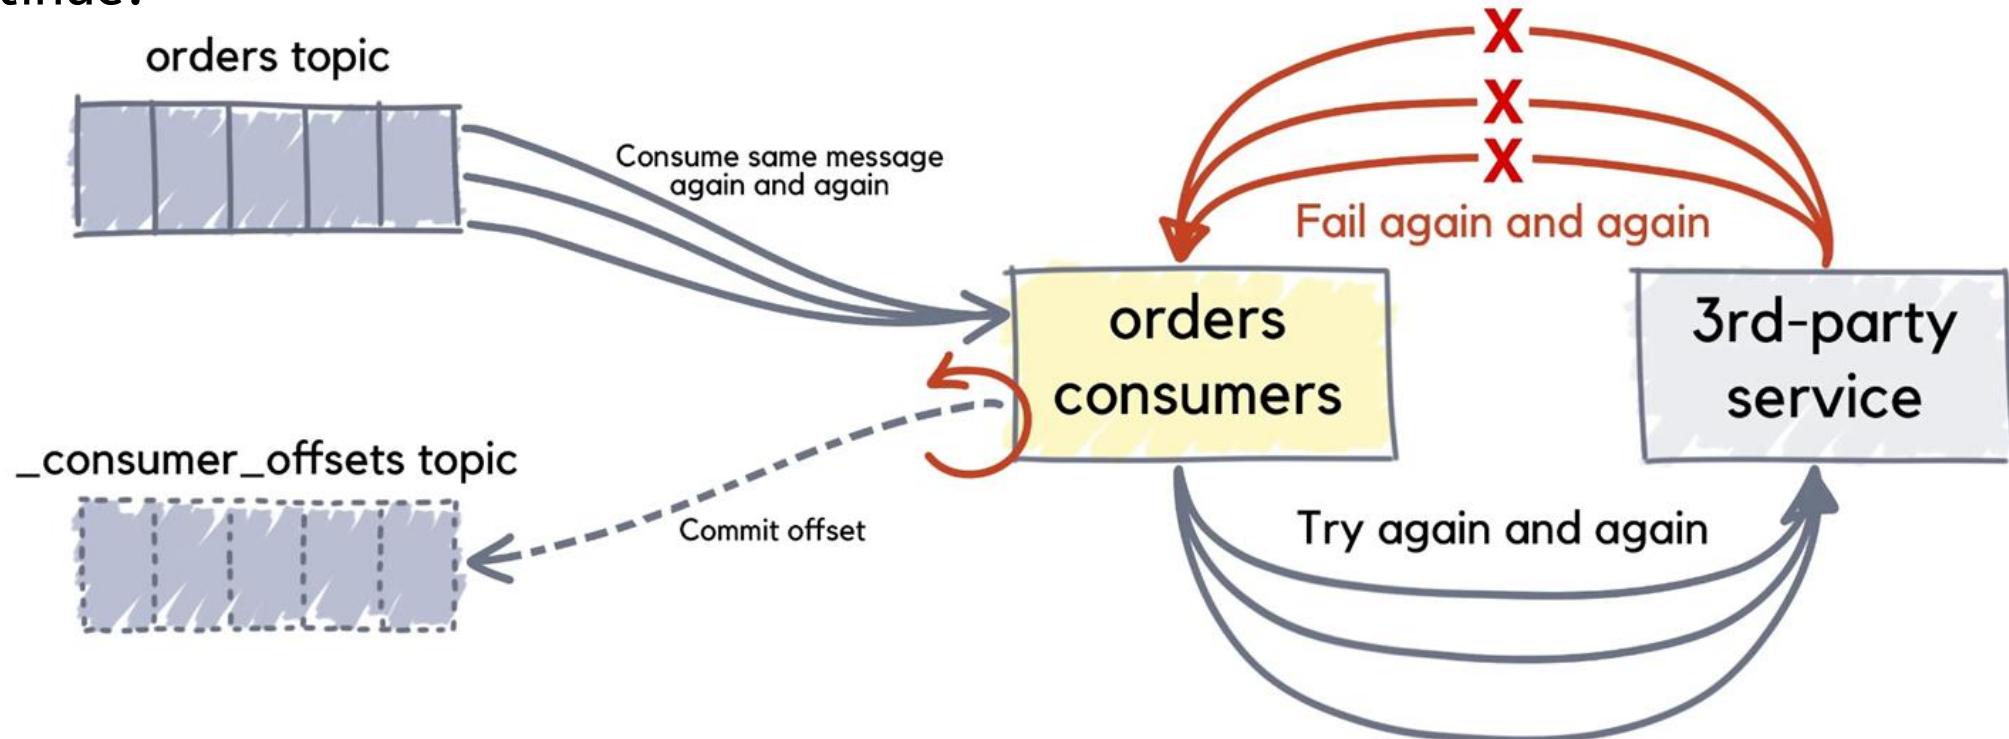
Dead-Letter Queue

Redpanda doesn't provide a built-in DLQ but lets you configure any topic as a custom dead-letter queue.



Retry handling

The best way to handle transient errors is to retry the operation several times, using fixed or incremental backoff intervals. If all retries fail, the message should be redirected to the DLT to prevent blocking and allow processing to continue.



Observability patterns post-MSK lessons

Metrics first

Track lag, throughput, and latency on a unified dashboard to spot issues fast.

Structured logging

Enrich logs with topic/partition/offset and request IDs for quick correlation.

Tracing

Use distributed tracing to follow messages end-to-end and pinpoint bottlenecks.

ConsuaAlerts & SLOsmers

Tie alerts to SLOs (e.g., lag, error rate, p99 latency) to act before users feel pain.

Failure patterns

Detect retry storms, DLQ growth, and replica issues and trigger auto-remediation.



Redpanda Platform & Operations (DevOps/Platform Focus)

**Redpanda Console Deep Dive
(with Terraform)**



- Console features for management, monitoring, debugging
- Admin tasks: GUI topic setup/editing
- Consumer lag, group metrics, time-travel debugging
- Access control, security configuration
- Terraform usage to automate console/admin tasks

Redpanda Console features

Redpanda Console is a web UI for managing and monitoring Redpanda clusters, making it easier to handle streaming data, troubleshoot issues, and optimize workloads.

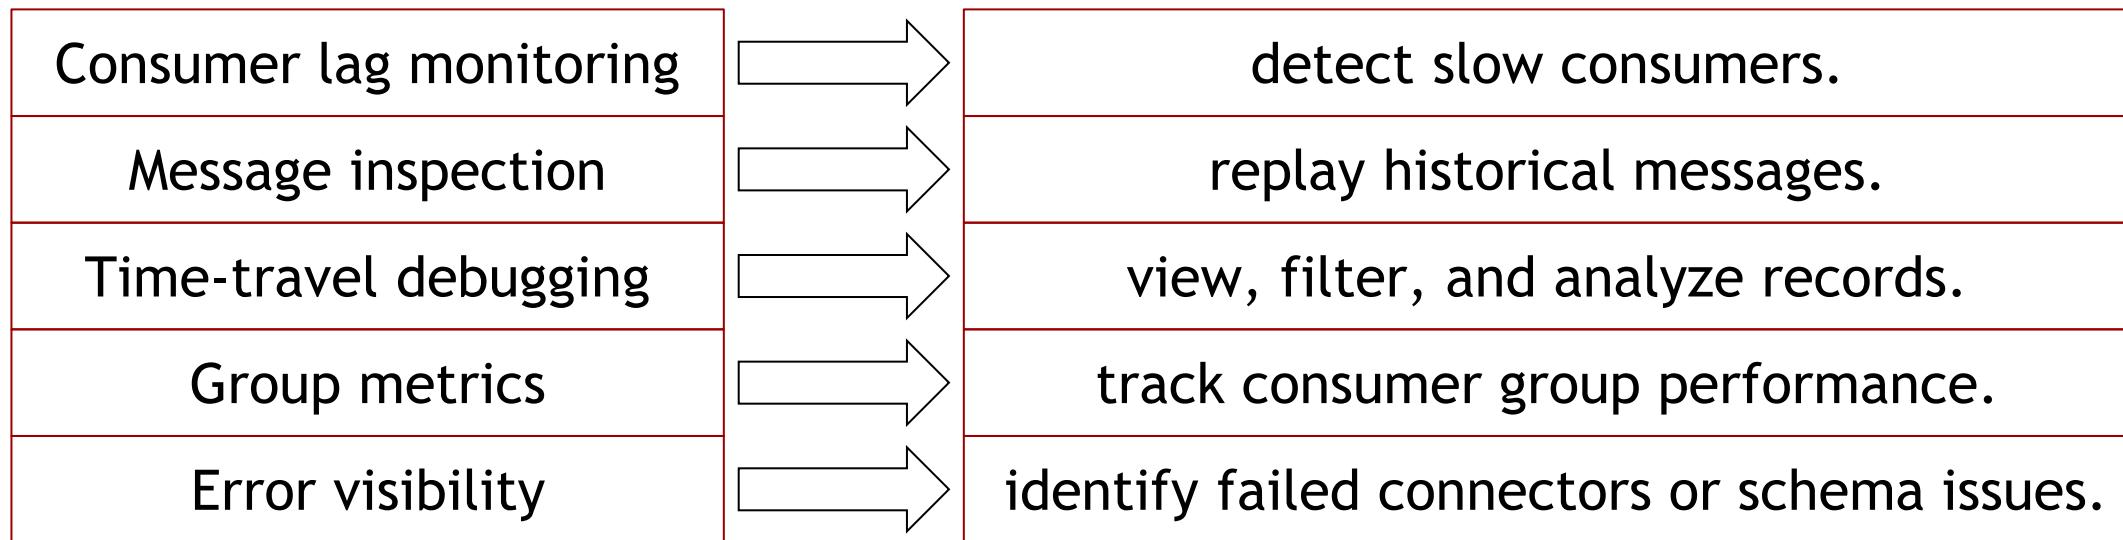
The screenshot shows the Redpanda Console interface. The left sidebar contains navigation links: Overview (selected), Topics, Schema Registry, Consumer Groups, Security, Quotas, Connect, Transforms, Reassign Partitions, superuser (with Preferences), and Collapse sidebar. The main content area has a header with 'Cluster' and tabs for 'Overview' (selected), 'Topics', and 'Schema Registry'. A message bar at the top right says 'Your Redpanda Enterprise trial will expire in 29 days. Request a full license.' Below this, cluster statistics are displayed: Running, 630 kiB, Redpanda v25.1.5, 3 of 3 brokers online, 7 topics, 93 replicas. The 'OVERVIEW' section includes 'BROKER DETAILS' (Broker IDs 0, 1, 2, all running) and 'CLUSTER DETAILS' (Services: Kafka Connect Not configured, Schema Registry Healthy, 1 schemas; Storage: Total Bytes 630 kiB, Primary 308 kiB, Replicated 322 kiB; Security: Service Accounts 1, ACLs 2). The bottom section, 'RESOURCES AND UPDATES', lists Documentation and CLI Tools.

Console features for management

- Create, edit, and delete topics
- Manage consumer groups and offsets
- Configure cluster settings
- Set up and manage access control
- Monitor and troubleshoot workloads

Debugging features

Redpanda Console provides a web-based interface that aids in debugging and troubleshooting Redpanda clusters and Kafka-compatible workloads.



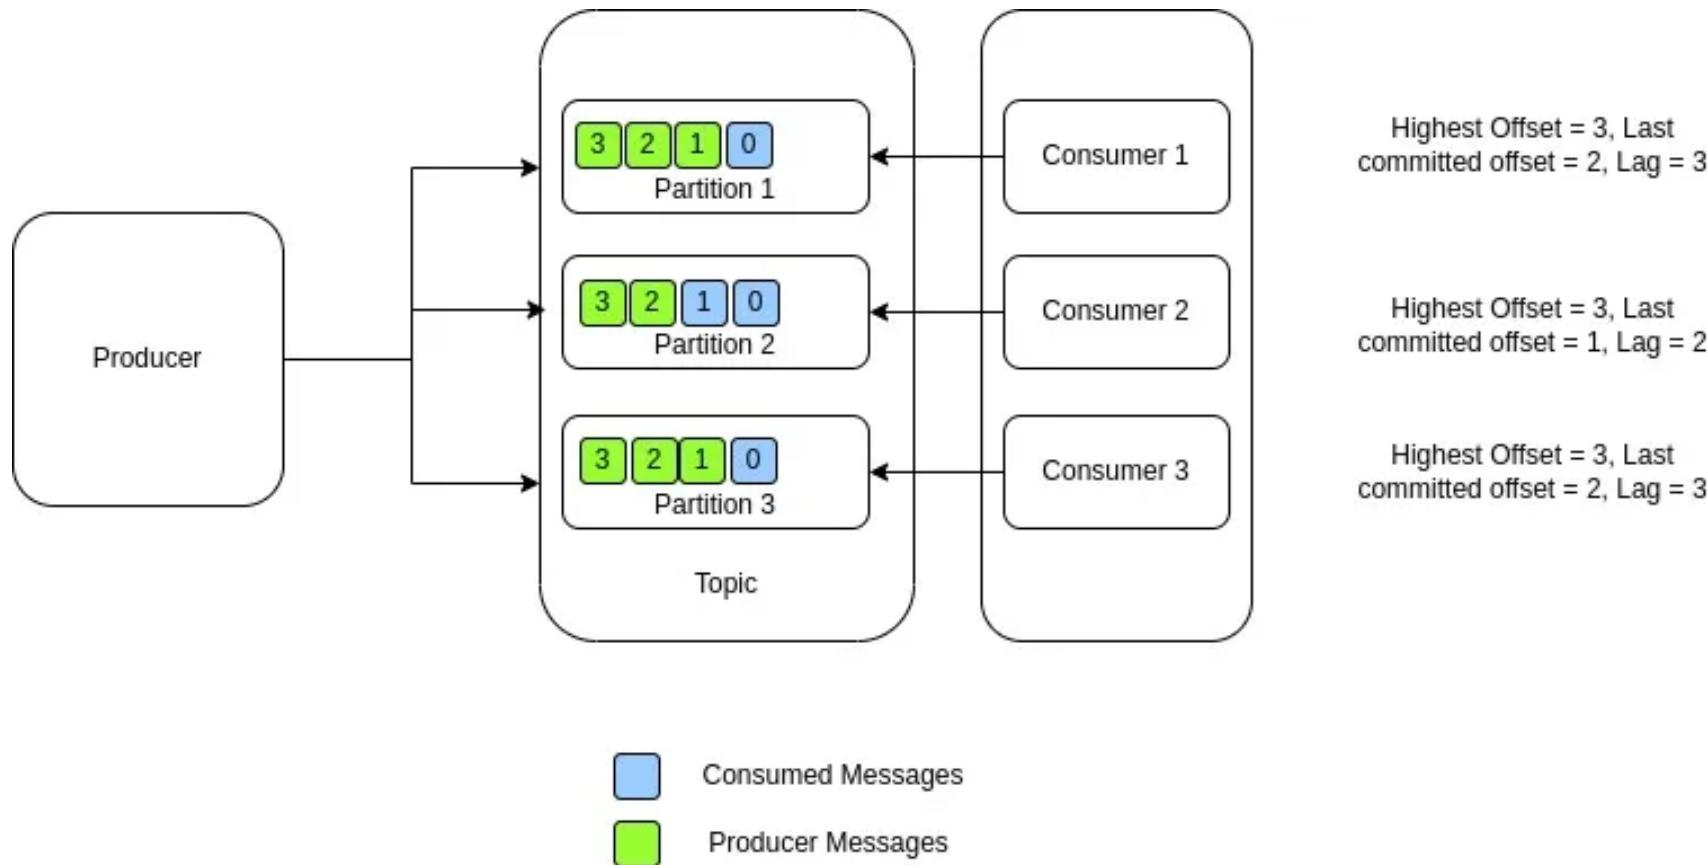
Admin Tasks: GUI Topic Setup & Editing

Redpanda Console provides an intuitive web-based interface that makes it easy for administrators to manage topics without relying on CLI commands. Through the GUI, you can:

- **Create Topics** - Set up new topics with partitions, replication, and advanced configurations.
- **Edit Topics** - Update topic settings, scale partitions, or adjust replication as needed.
- **Delete Topics** - Remove unused topics to keep the cluster clean and efficient.
- **Provides Visibility** - View real-time topic configurations at a glance.
- **Reduces Errors** - Minimize risks compared to manual CLI commands.
- **Simplifies Management** - Manage topics easily through a no-code GUI.

Consumer lag in Redpanda

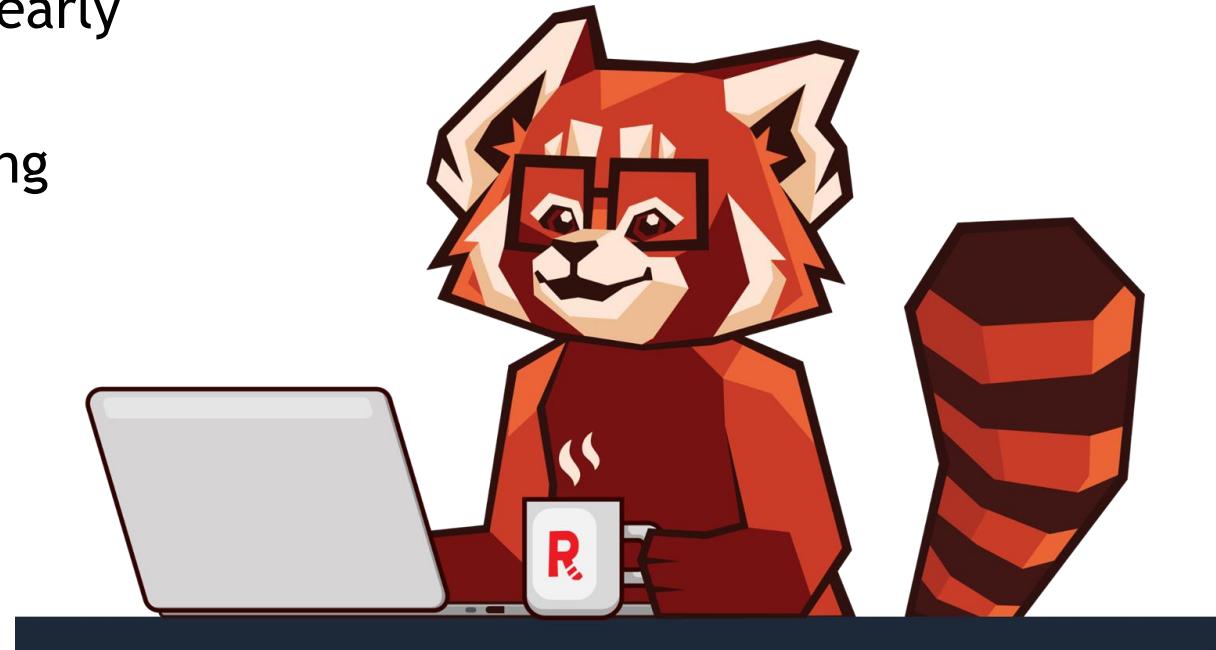
Consumer lag is the difference between the last offset stored by the broker and the last committed offset for that partition.



Group Metrics in Redpanda

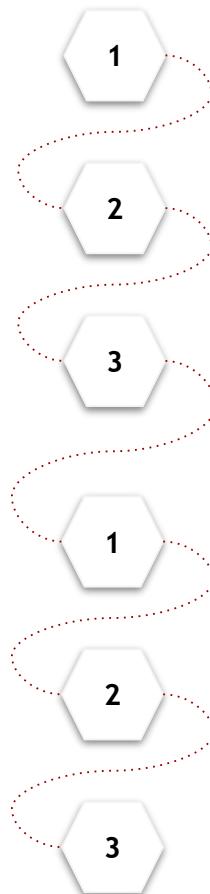
Why Monitor Consumer Groups?

- 1 Ensure consumers keep up with producers
- 2 Detect performance bottlenecks early
- 3 Maintain healthy stream processing



Group Metrics in Redpanda

Core Metrics



redpanda_kafka_consumer_group_committed_offset

Last committed offset for a group on a partition

redpanda_kafka_consumer_group_consumers

Number of active consumers in the group

redpanda_kafka_consumer_group_topics

Number of topics subscribed to by the group

redpanda_kafka_group_lag_max

Maximum lag across any partition in the group

redpanda_kafka_group_lag_sum

Total lag across all partitions in the group

redpanda_kafka_max_offset

Latest broker offset for a partition (used to calculate lag)



Time-travel debugging

Time-travel debugging in Redpanda is a feature primarily provided through the Redpanda Console.

- **Explore historical messages** - Navigate past messages to see data at specific points in time.
- **Identify & debug problems** - Pinpoint messages or events that caused errors or unexpected behavior.
- **Gain visibility** - Visualize and analyze data streams interactively to spot anomalies.



Access control

Redpanda uses Access Control Lists (ACLs) and Role-Based Access Control (RBAC) to manage user permissions and access to resources.

Access Control Lists (ACLs)

Access control lists (ACLs) provide a way to configure fine-grained access to provisioned users. ACLs work with SASL/SCRAM and with mTLS with principal mapping for authentication.

Role-Based Access Control (RBAC)

Role-based access control (RBAC) lets you manage user permissions at scale with a simple interface, working alongside all supported authentication methods. It ensures secure, consistent, and centralized access management across resources.

Configure Authentication

Authentication verifies users and apps connecting to Redpanda, with APIs supporting multiple methods and configurable listeners.

API	Supported Authentication Methods
Kafka API	<ul style="list-style-type: none">• SASL<ul style="list-style-type: none">◦ SASL/SCRAM◦ SASL/PLAIN◦ SASL/OAUTHBEARER (OIDC)◦ SASL/GSSAPI (Kerberos)• mTLS
Admin API	<ul style="list-style-type: none">• Basic authentication• OIDC
HTTP Proxy (PandaProxy)	<ul style="list-style-type: none">• Basic authentication• OIDC
Schema Registry	<ul style="list-style-type: none">• Basic authentication• OIDC

Terraform for Redpanda Console/Admin Automation

Why Terraform?

Codify Redpanda configuration as **Infrastructure as Code (IaC)**

Ensure **consistency and repeatability** across environments

Enable **version control** and easy rollbacks via Git

Automate **cluster + console tasks** with a single `terraform apply`

Terraform for Redpanda Console/Admin Automation

What Can Be Automated?

Provision and configure **topics** (partitions, replication, retention)

Apply **ACLs** and **RBAC policies** for secure access

Manage **connectors, schemas, and consumer groups**

Deploy **multi-environment setups** (Dev / Stage / Prod)



Redpanda Platform & Operations (DevOps/Platform Focus)

Local Development with Docker



- Running Redpanda and Console via Docker/Docker Compose
- Producing/consuming test messages locally
- Local connector testing and debugging pipelines

Docker/Docker Compose & Redpanda



- Simplifies deployment with containers
- No manual installation required
- Easy to scale and test locally
- Consistent dev/test environments

Steps to Run

1. Install Docker & Docker Compose

- Ensure Docker engine is running

2. Create a docker-compose.yml

- Define Redpanda broker service
- Define Redpanda Console service

3. Start Services

- `docker-compose up -d`

4. Access Redpanda Console

- Open <http://localhost:8080> in browser

5. Manage Topics & Monitor Metrics

- Use Console UI for topics, consumers, ACLs, connectors

With just one file + `docker-compose up`, you get a fully functional Redpanda cluster + Console running locally.

Use Cases for Local Setup

- **Developer testing:**

Quickly test producers/consumers without cloud infra

- **Demo/Training environments:**

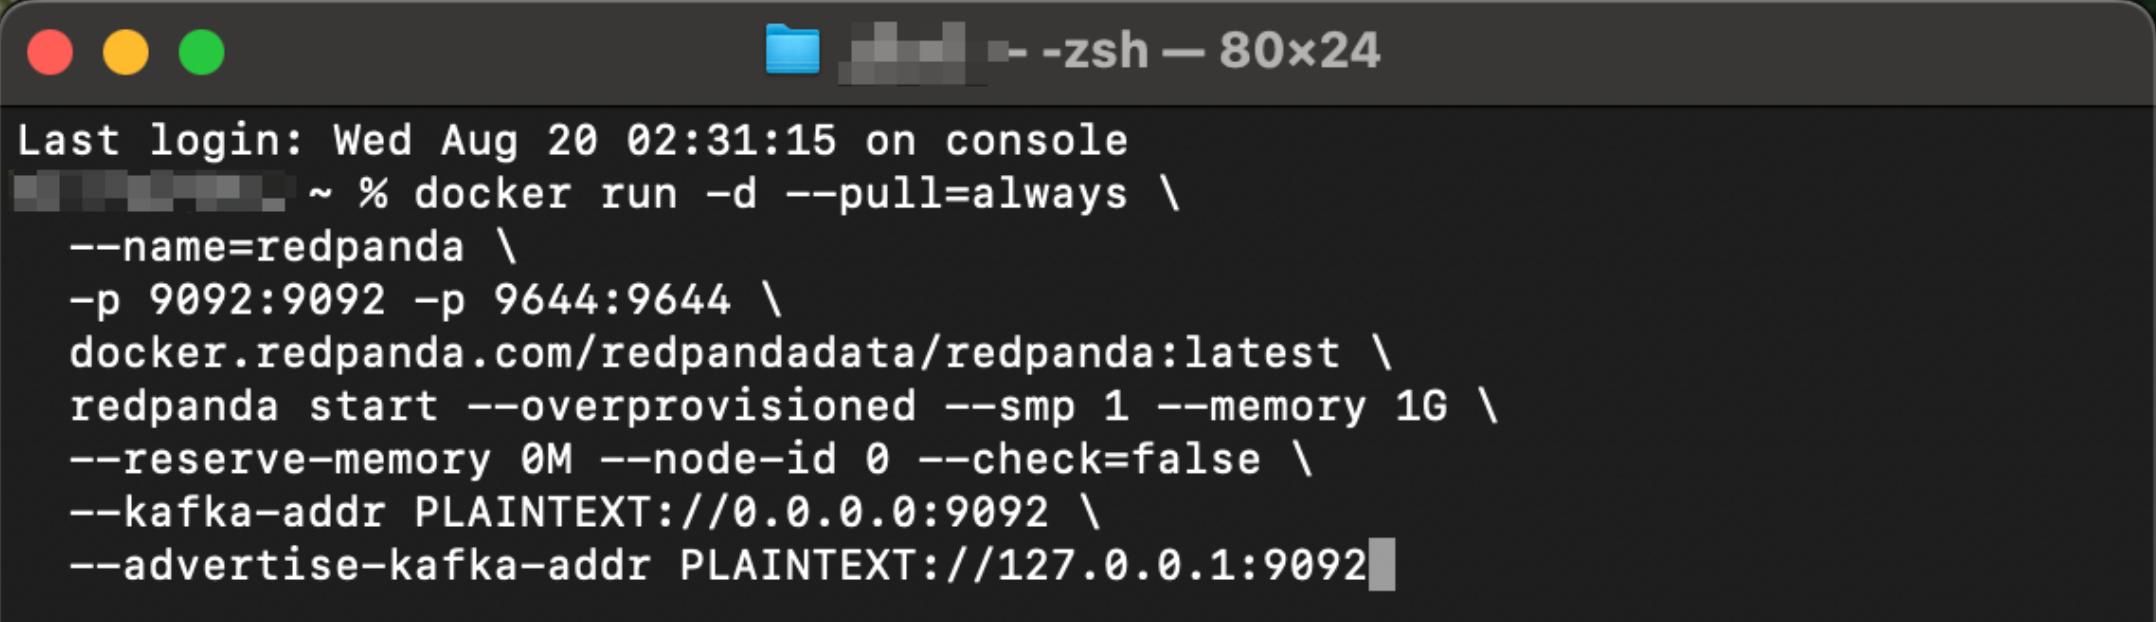
Spin up Redpanda + Console in seconds

- **CI/CD pipelines:**

Validate Kafka apps against a real broker in containerized tests

Producing/consuming test messages locally

Step 1: Start Redpanda



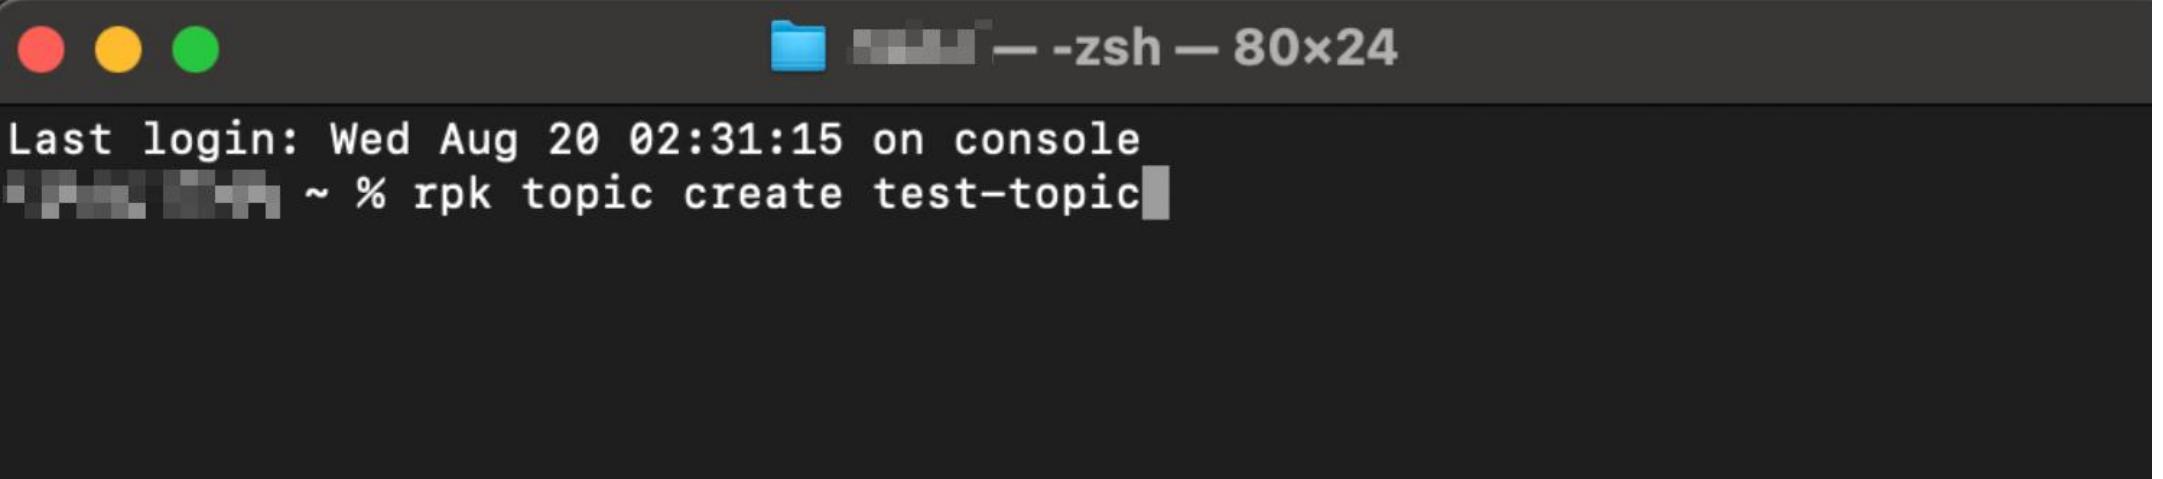
The screenshot shows a terminal window with a dark background and light-colored text. At the top, there are three colored circles (red, yellow, green) followed by a blue folder icon and the text "- -zsh - 80x24". The terminal output starts with "Last login: Wed Aug 20 02:31:15 on console" and then shows a command being typed:

```
Last login: Wed Aug 20 02:31:15 on console
~ % docker run -d --pull=always \
--name=redpanda \
-p 9092:9092 -p 9644:9644 \
docker.redpanda.com/redpandadata/redpanda:latest \
redpanda start --overprovisioned --smp 1 --memory 1G \
--reserve-memory 0M --node-id 0 --check=false \
--kafka-addr PLAINTEXT://0.0.0.0:9092 \
--advertise-kafka-addr PLAINTEXT://127.0.0.1:9092
```

This command runs a **single-node Redpanda cluster** in Docker with 1 CPU, 1GB RAM, exposing **9092 (Kafka)** and **9644 (Admin)** for easy local testing.

Producing/consuming test messages locally

Step 2: Create a Topic

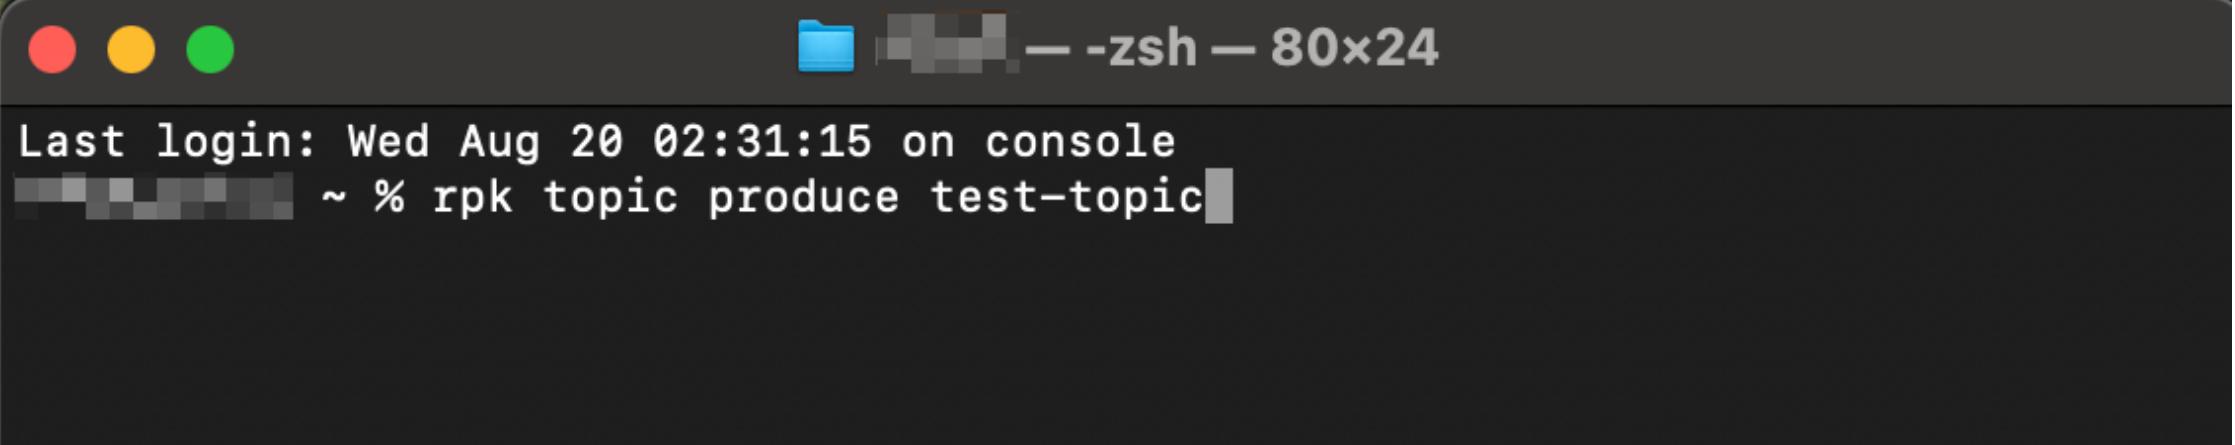


A screenshot of a terminal window titled "-zsh - 80x24". The window shows a user's last login information: "Last login: Wed Aug 20 02:31:15 on console". Below this, the command "rpk topic create test-topic" is typed into the terminal. The terminal has a dark background with light-colored text and icons.

- `rpk` is the Redpanda CLI tool (Redpanda Keeper).
- `topic create` is the command to create a new topic in Redpanda.
- `test-topic` is the name of the topic being created, which will be used later for producing (sending) and consuming (reading) test messages.

Producing/consuming test messages locally

Step 3: Produce Messages

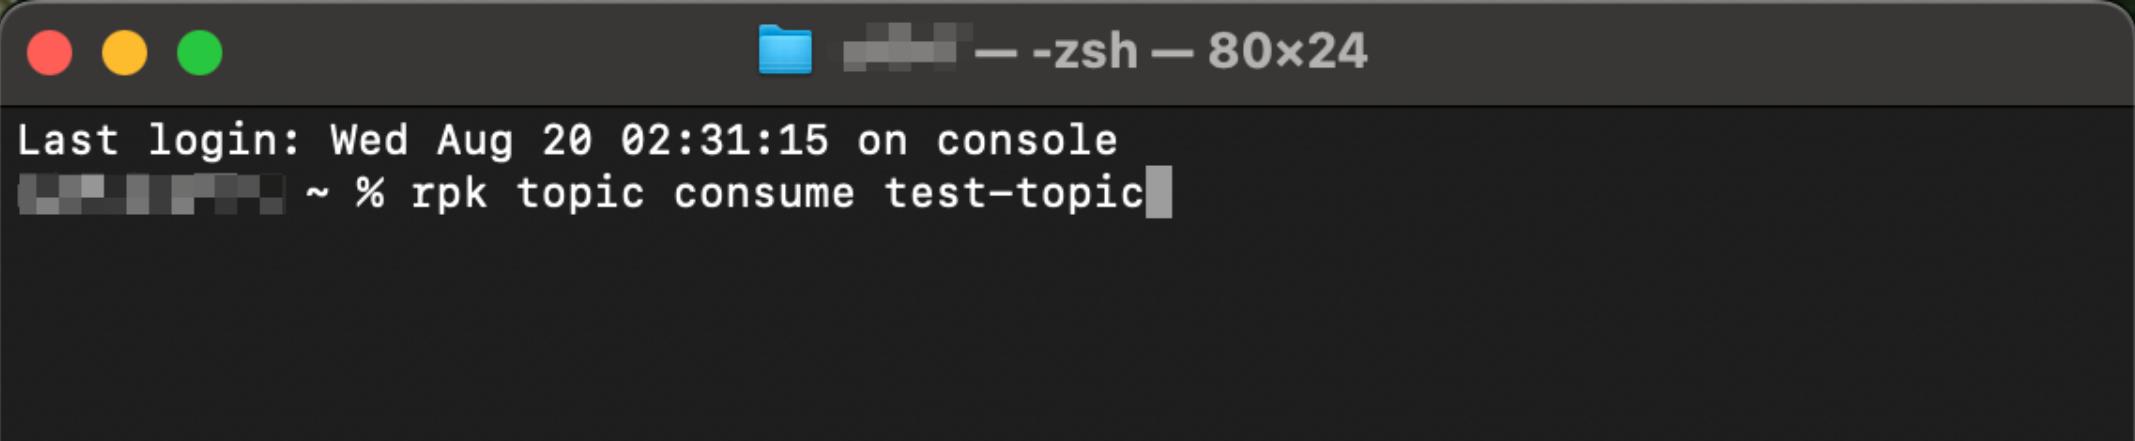


A screenshot of a terminal window titled '-zsh - 80x24'. The window shows a user's last login information: 'Last login: Wed Aug 20 02:31:15 on console'. Below this, the command 'rpk topic produce test-topic' is typed into the terminal. The background of the terminal window is dark, and the text is white.

- **rpk** → Redpanda's CLI tool (similar to Kafka's [kafka-console-producer.sh](#)).
- **topic produce** → Command to produce (send) messages into a topic.
- **test-topic** → The topic name where the messages will be written.

Producing/consuming test messages locally

Step 4: Consume Messages



A screenshot of a macOS terminal window titled "-zsh - 80x24". The window shows the following text:

```
Last login: Wed Aug 20 02:31:15 on console  
[REDACTED] ~ % rpk topic consume test-topic
```

- `rpk` → The Redpanda CLI tool.
- `topic` → Subcommand to manage topics.
- `consume` → Starts a consumer that listens for messages.
- `test-topic` → The name of the topic to consume from.

Why Local Testing?

- 1 Safe environment to test connectors before production
- 2 Faster debugging cycle with minimal setup
- 3 Identify schema, serialization, and transformation issues early
- 4 Reduce risk of failures in live systems

Local Setup

Step: 1

Run Redpanda + Connect + Console via Docker Compose

Step: 2

Place connector JARs in `/var/lib/kafka-connect/plugins`

Step: 3

Create **test topics**: `test-input`, `test-output`

Step: 4

Use `rpk` CLI to produce and consume test events



THANK YOU